香港科技大学(广州)
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY (GUANGZHOU)



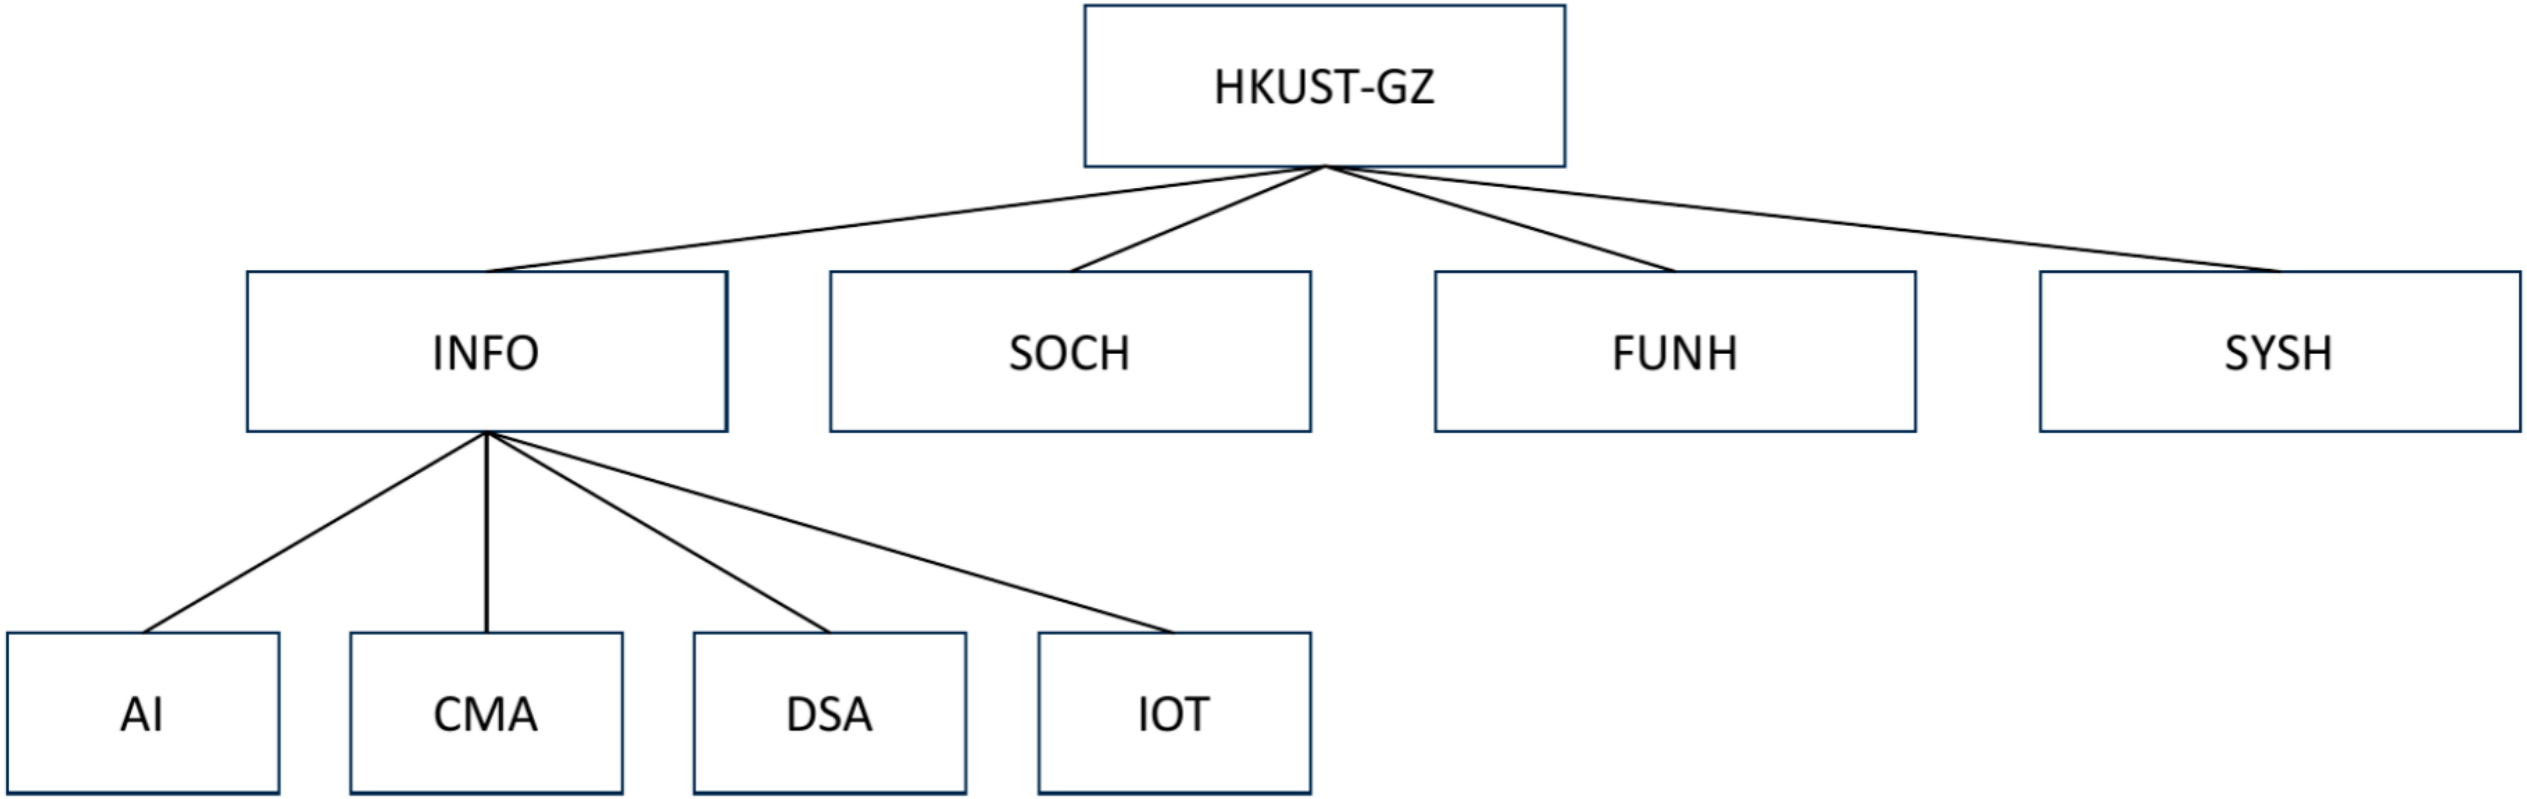
Algorithms on Trees

- Tree Terminology
- Tree Traversal
- Tree DP

Yanlin Zhang & Shangqi Lu | DSAA 2043 Fall 2026

Tree Terminology

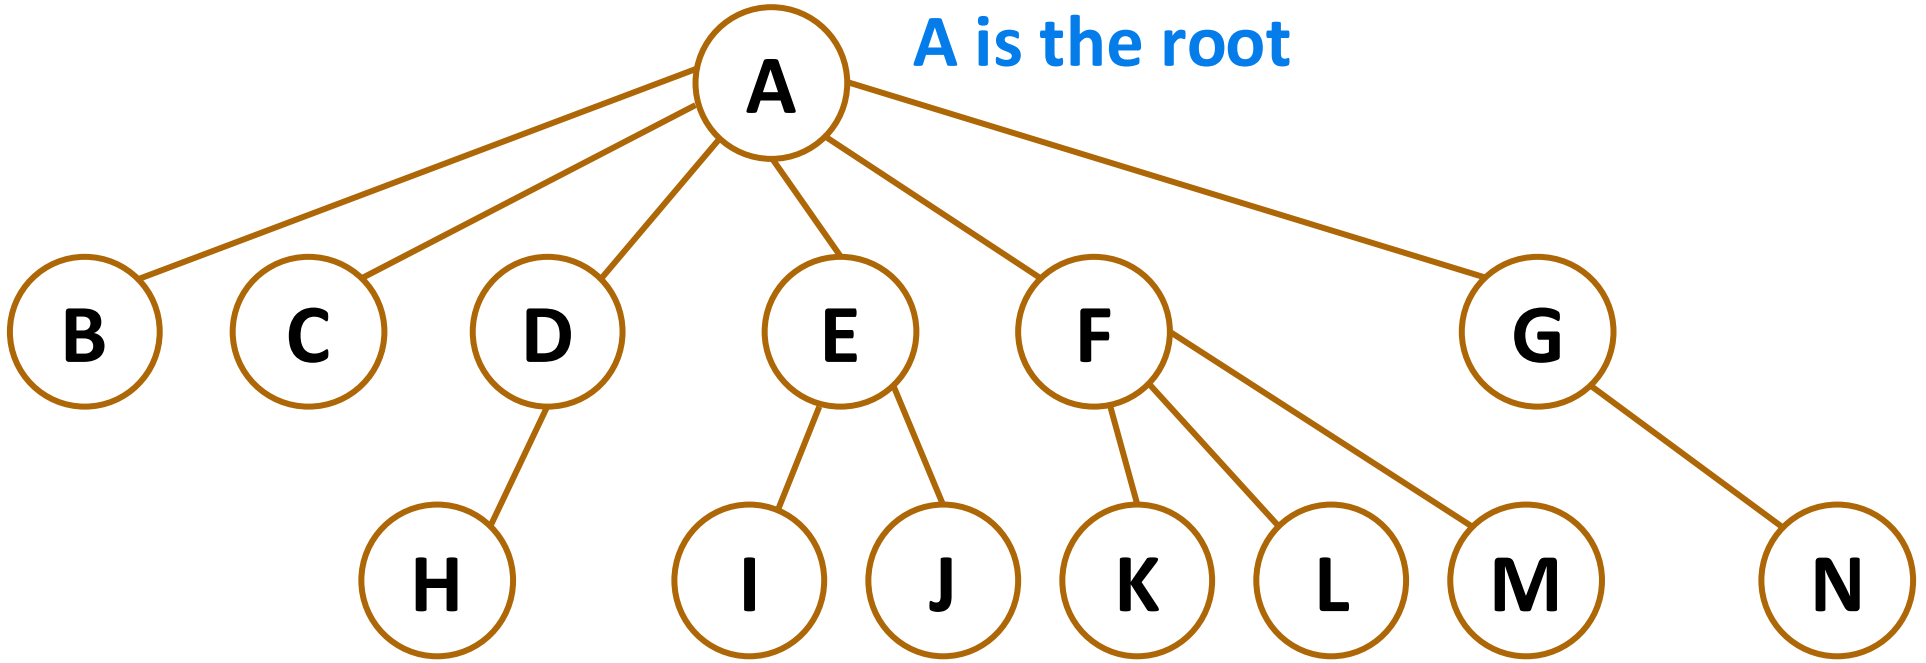
Trees Can Describe Hierarchies



Trees are inherently recursive

What is a Tree (in Computer Science)?

- A tree is a collection of **nodes**, which can be empty. If it is not empty, there is a "root" node,***r***, and ***zero or more non-empty subtrees***, ***T1, T2, ..., Tk***, whose roots are connected by adirected edge from ***r***.

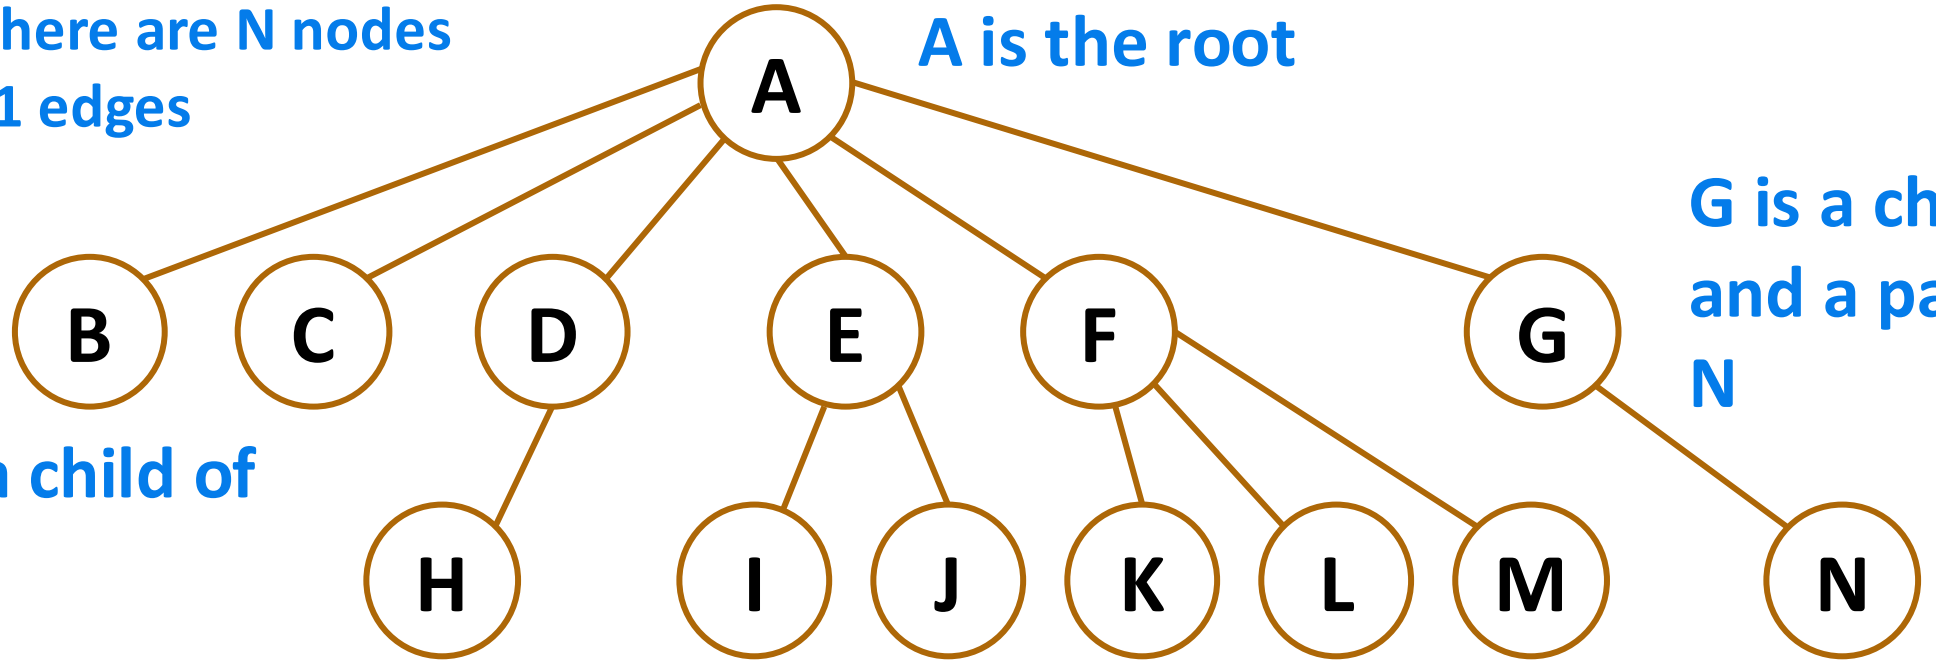


Tree Terminology

What is a Tree (in Computer Science)?

- A tree is a collection of **nodes**, which can be empty. If it is not empty, there is a "root" node,***r***, and ***zero or more non-empty subtrees***, ***T1, T2, ..., Tk***, whose roots are connected by adirected edge from ***r***.

Note: there are N nodes and N-1 edges



A is the root

G is a child of A and a parent of N

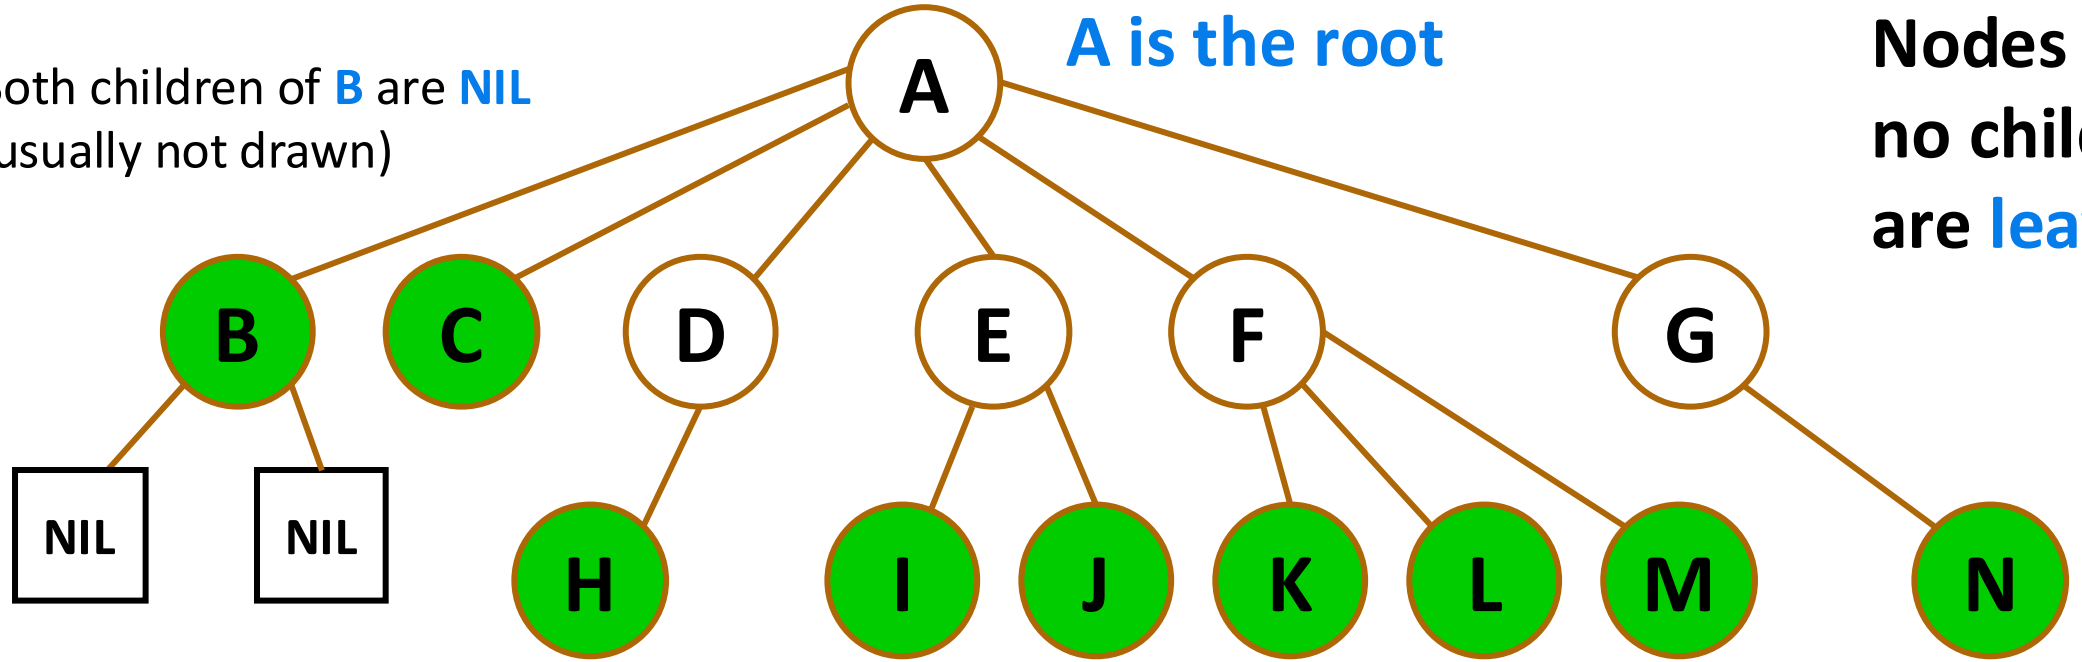
B is a child of A

Tree Terminology

What is a Tree (in Computer Science)?

- A tree is a collection of **nodes**, which can be empty. If it is not empty, there is a "root" node,***r***, and ***zero or more non-empty subtrees***, ***T1, T2, ..., Tk***, whose roots are connected by adirected edge from ***r***.

Both children of **B** are **NIL** (usually not drawn)

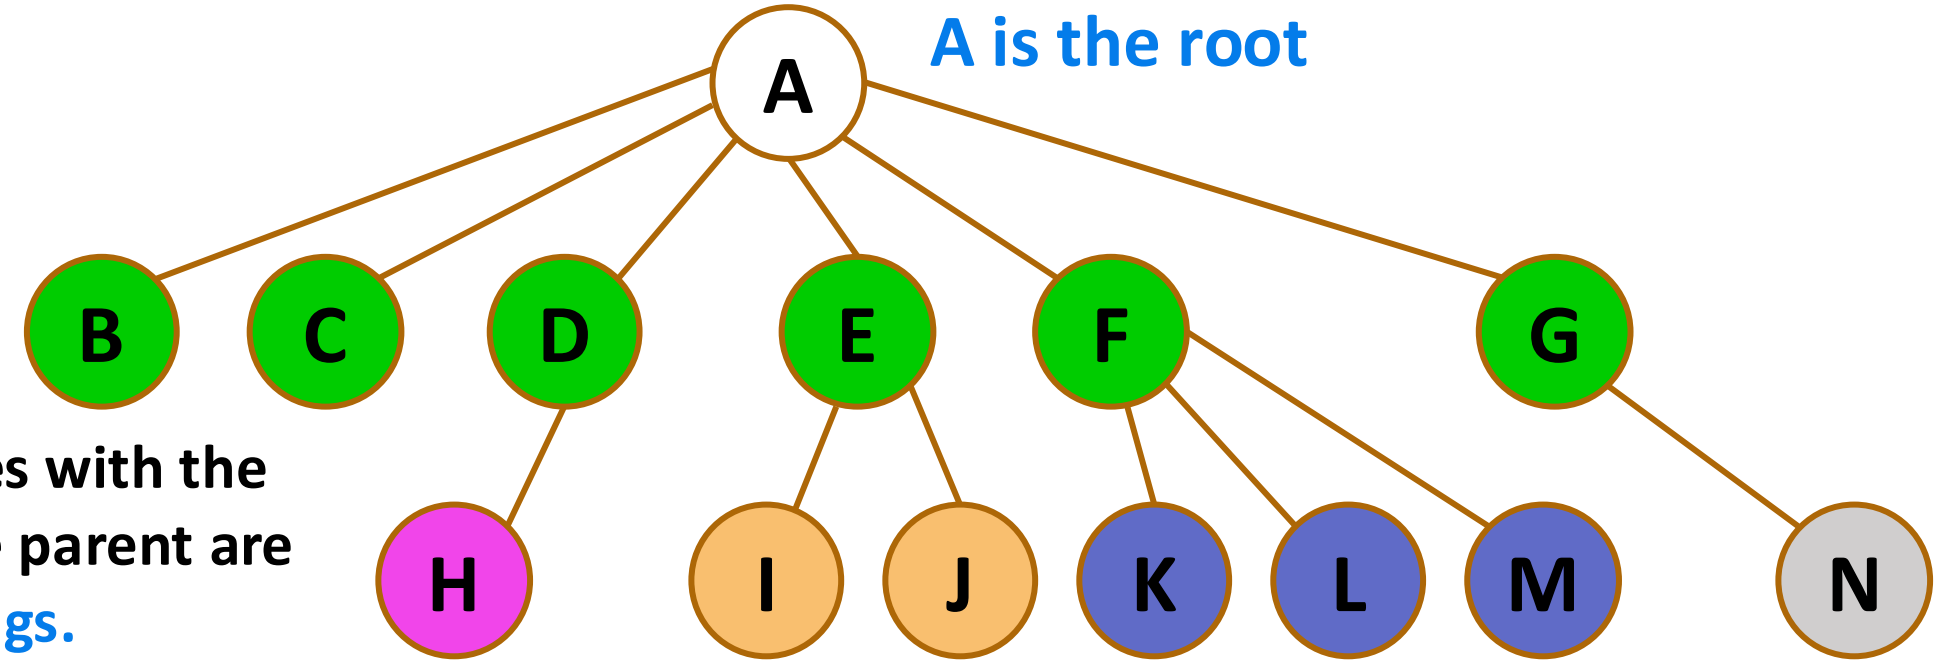


Nodes with no children are leaves

Tree Terminology

What is a Tree (in Computer Science)?

- A tree is a collection of **nodes**, which can be empty. If it is not empty, there is a "root" node,***r***, and ***zero or more non-empty subtrees***, ***T1, T2, ..., Tk***, whose roots are connected by adirected edge from ***r***.

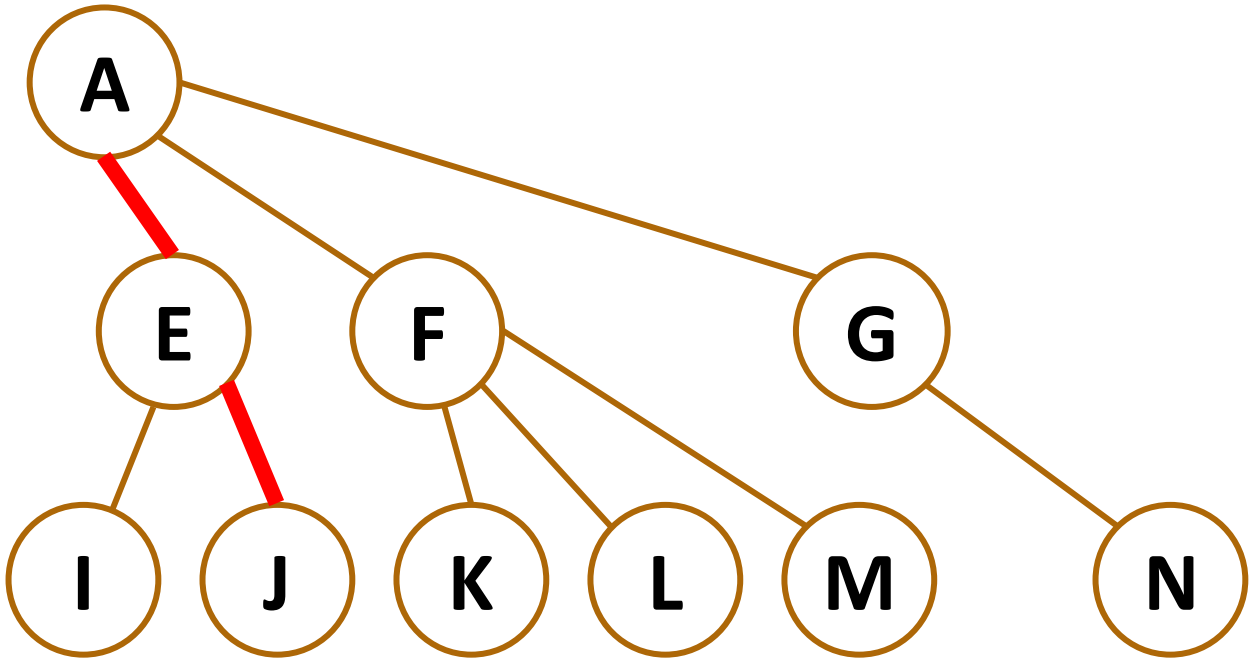


Nodes with the same parent are siblings.

Tree Terminology

We can define a path from a parent to its children.

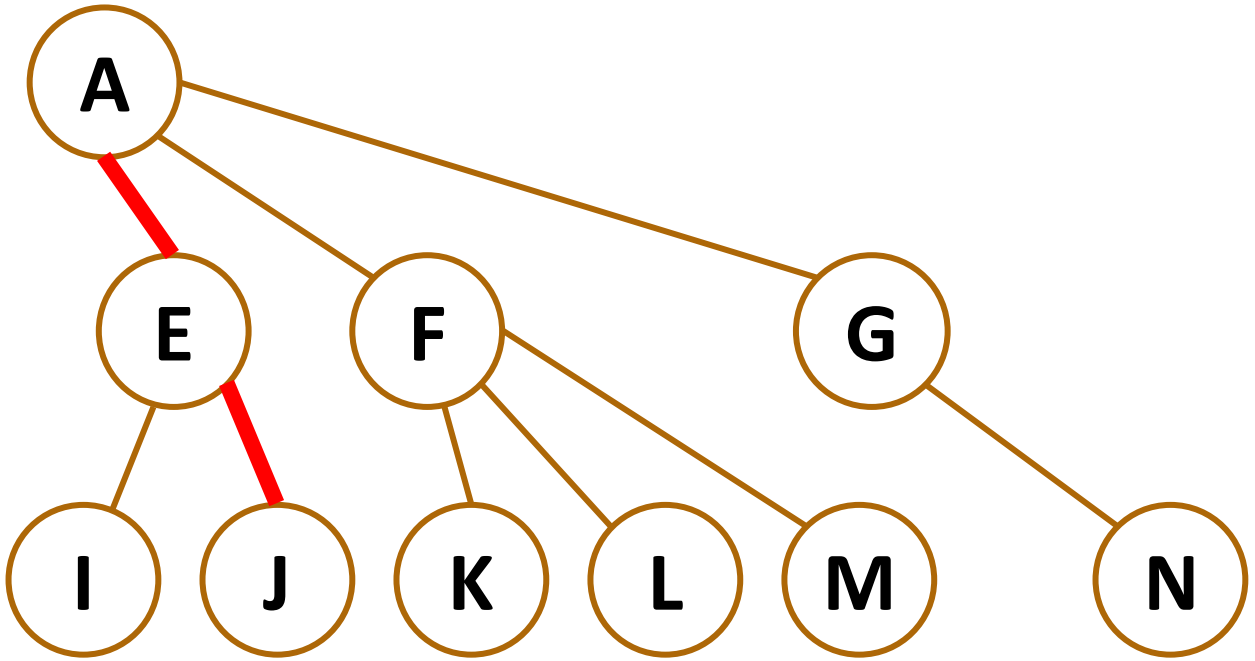
The path A-E-J has a length of two(the number of edges)



Tree Terminology

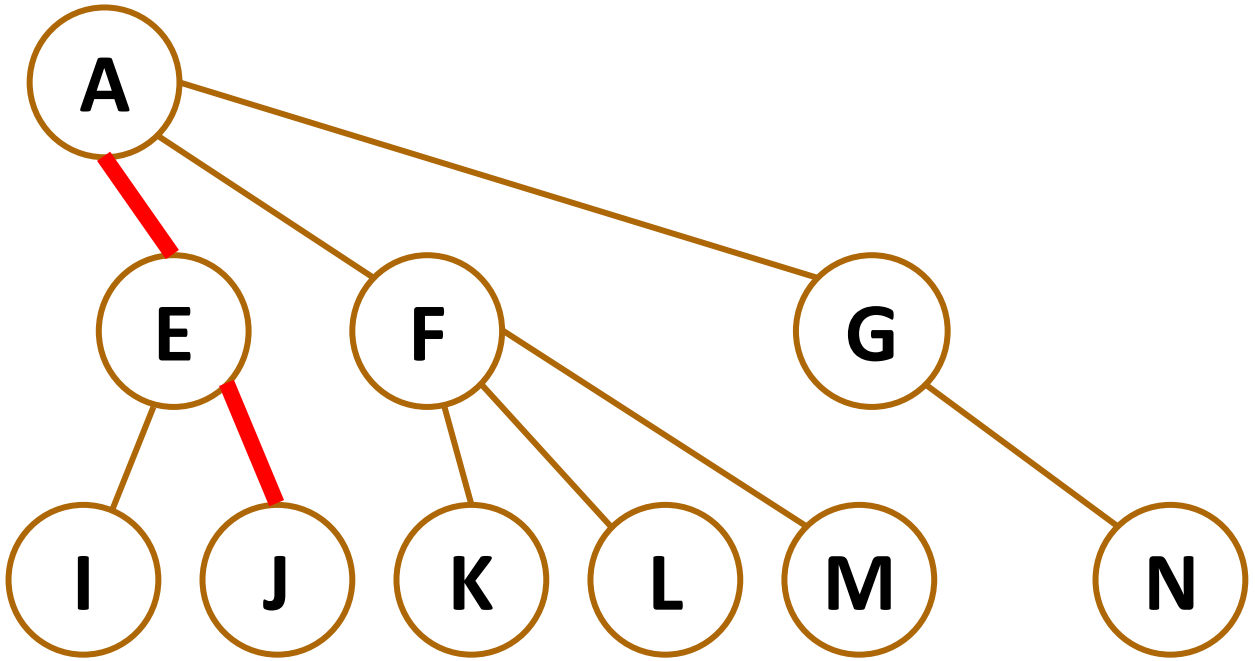
The depth of a node is the length from the root. The depth of node J is 2. The depth of the root is 0.

The height of a node is the longest path from the node to a leaf. The height of node F is 1. The height of all leaves is 0.



Tree Terminology

The height of a tree is the height of the root (in this case,the height of the tree is 2).

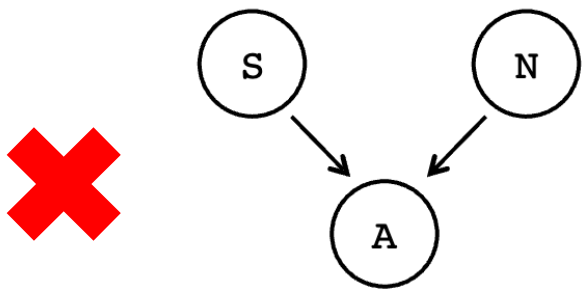


Tree Terminology

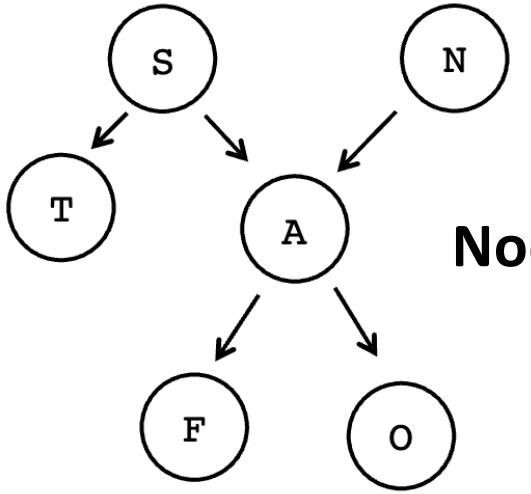
Trees can have only one parent, and cannot have cycles







Node A has two parents



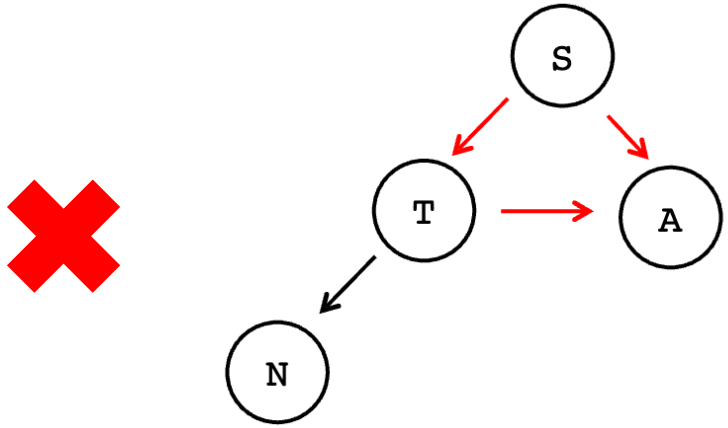
Node A has two parents

Tree Terminology

Trees can have only one parent, and cannot have cycles







not a tree: the red edges make a cycle

Binary Tree Terminology

A tree is a binary tree so that

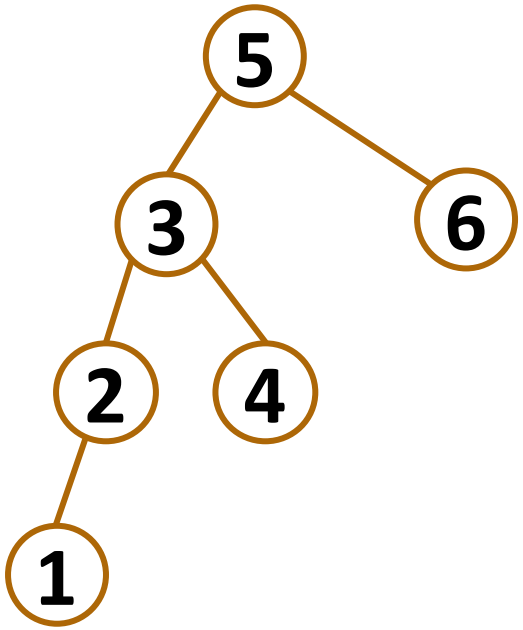
- Each node has at most two children

A **Binary Search Tree (BST)** is a binary tree so that

- Every LEFT descendant of a node has key less than that node.
- Every RIGHT descendant of a node has key larger than that node.

2 is a descendant of 5.

5 is an ancestor of 2.



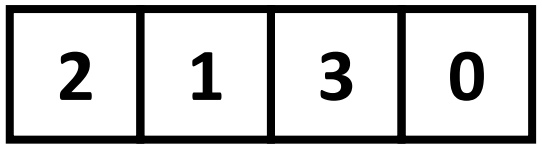
Traversal

Iterate through data structure

Iterate: process elements of data structure

- Sum all elements
- Print each element
- ...

Array



Forwards: 2, 1, 3, 0

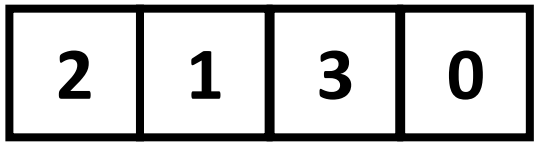
Backwards: 0, 3, 1, 2

Iterate through data structure

Iterate: process elements of data structure

- Sum all elements
- Print each element
- ...

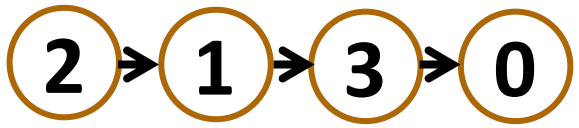
Array



Forwards: 2, 1, 3, 0

Backwards: 0, 3, 1, 2

Linked List



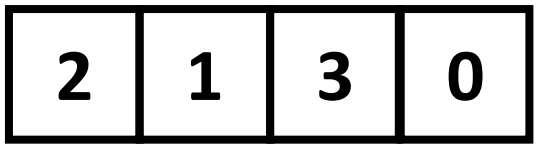
Forwards: 2, 1, 3, 0

Iterate through data structure

Iterate: process elements of data structure

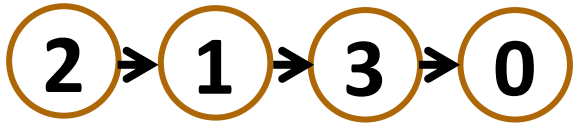
- Sum all elements
- Print each element
- …

Array



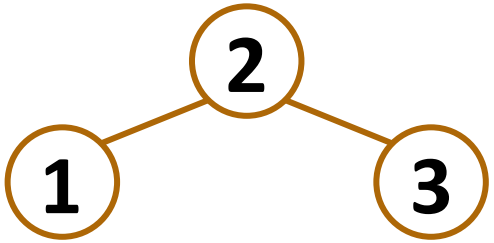
Forwards: 2, 1, 3, 0
Backwards: 0, 3, 1, 2

Linked List



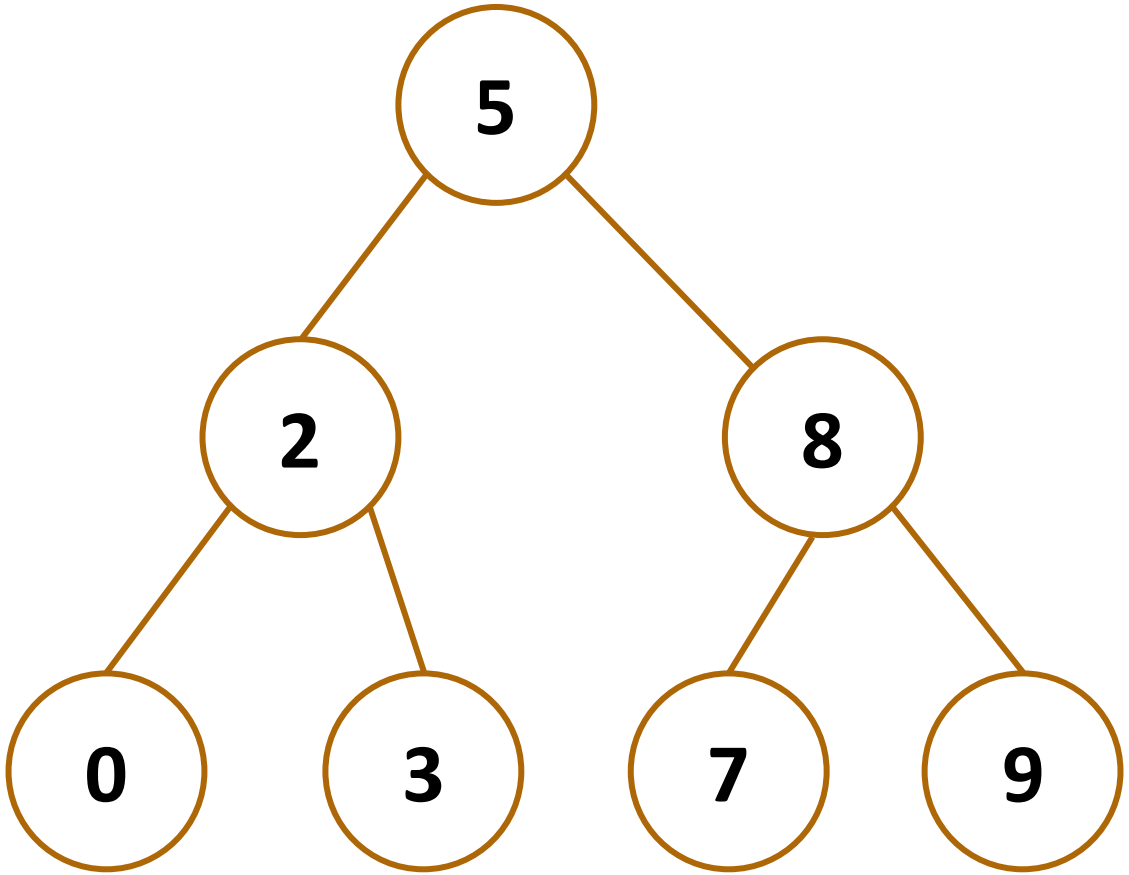
Forwards: 2, 1, 3, 0

BST



???

Iterate through data structure



Discuss: What would a reasonable order be?

Tree traversals

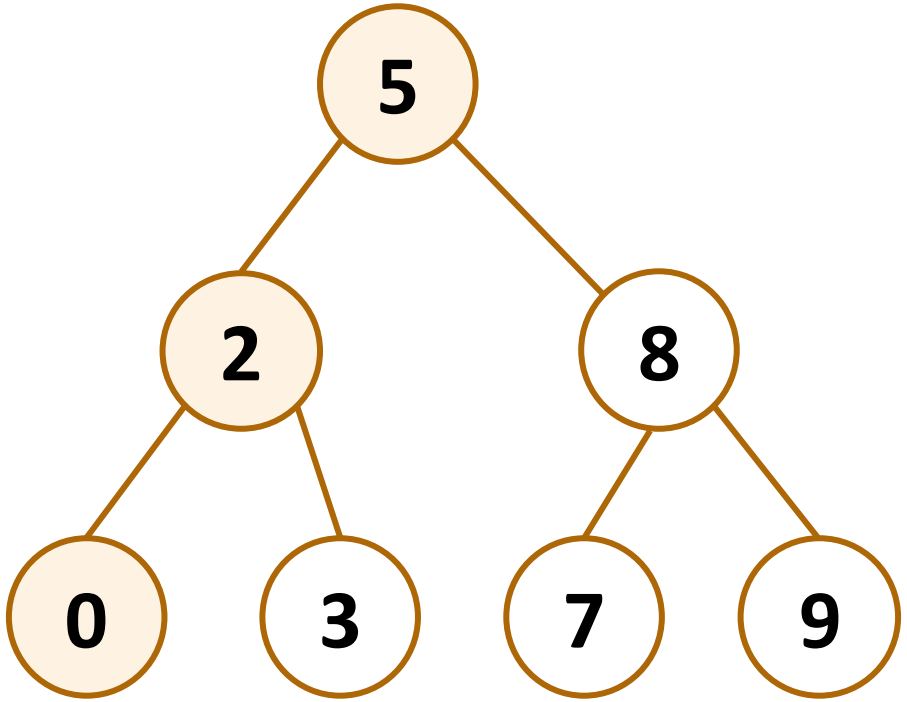
Iterating through tree is **tree traversal : A systematic way of accessing (or visiting) all the nodes of a tree**

Well-known recursive tree traversal algorithms:

- **Depth First Search (DFS)**
  - ✓ Preorder
  - ✓ Inorder
  - ✓ Postorder
- **Breadth First Search (BFS)**
  - ✓ Level order

DFS: go deep, then backtrack

- Start at the root.
- Keep following one child until you reach a leaf.
- Then backtrack to the most recent node with an unvisited child.



DFS: go deep, then backtrack

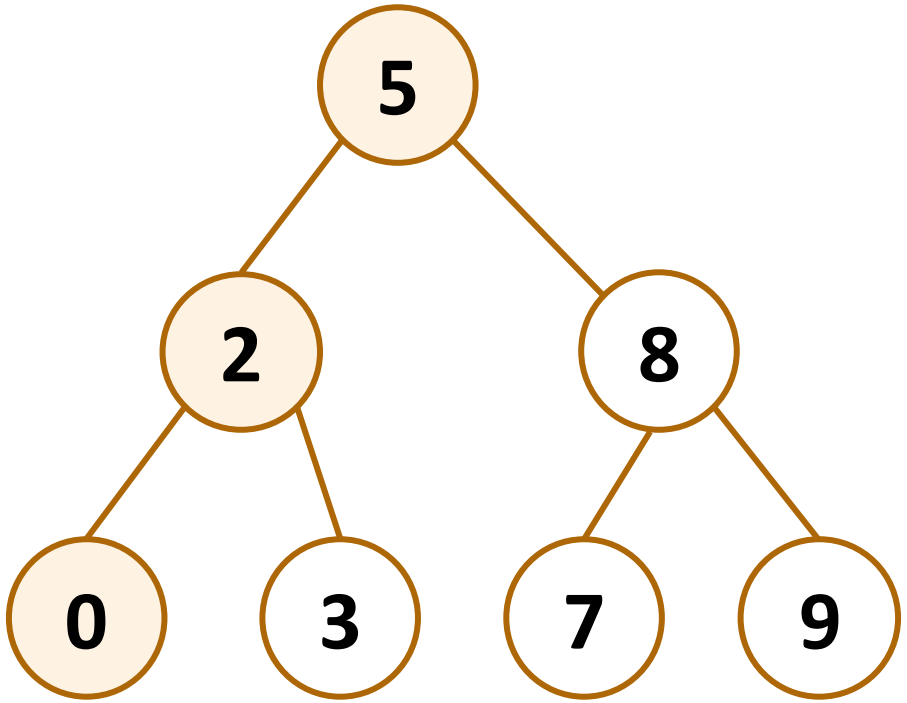
香港科技大学(广州)
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY (GUANGZHOU)

Recursive idea

```
1 def dfs(u):
2     visit(u)
3     for v in children[u]:
4         dfs(v)
```

visit u first

then recursively process every child

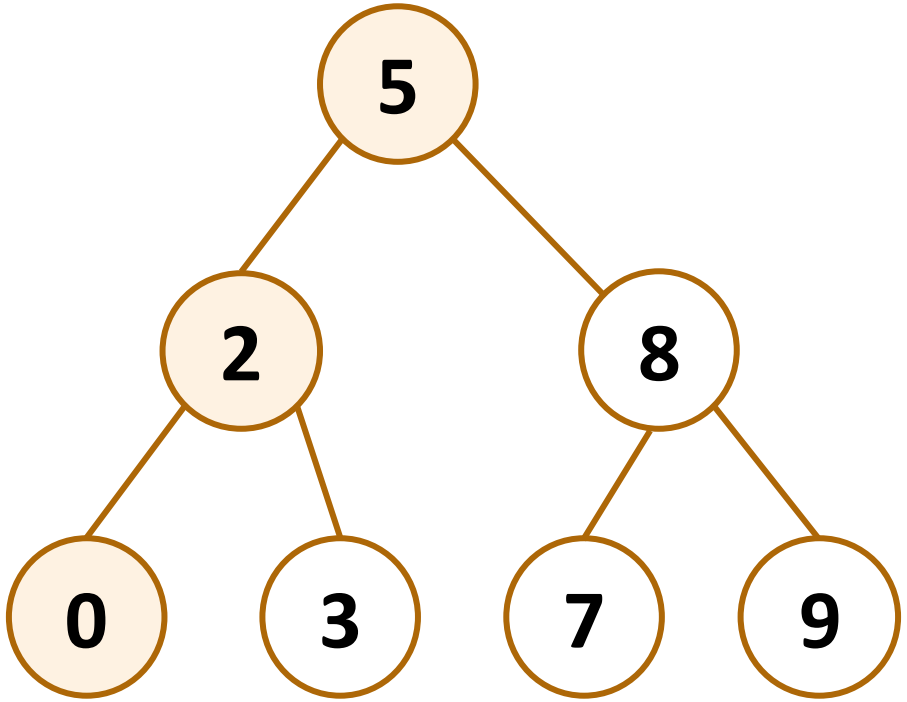


DFS walkthrough

Visited so far

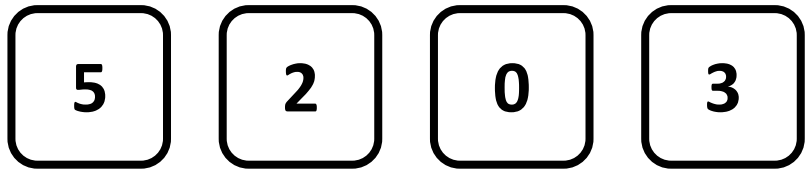


We reached a leaf, so the next action is backtracking.

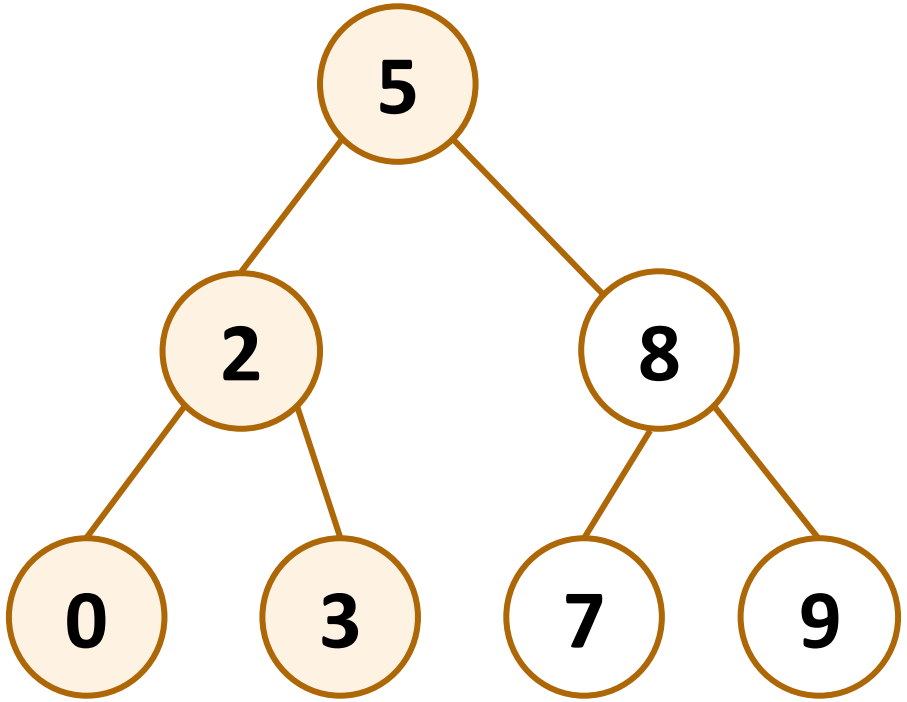


DFS walkthrough

Visited so far

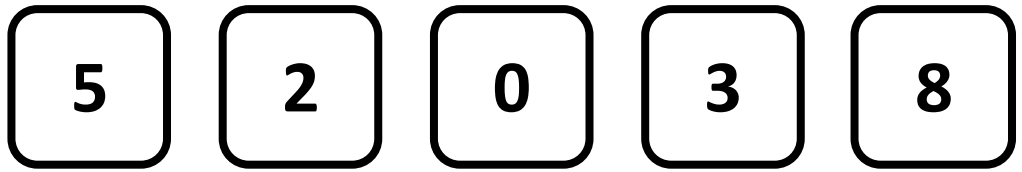


Backtrack to **2** and then continue with its next child **3**.

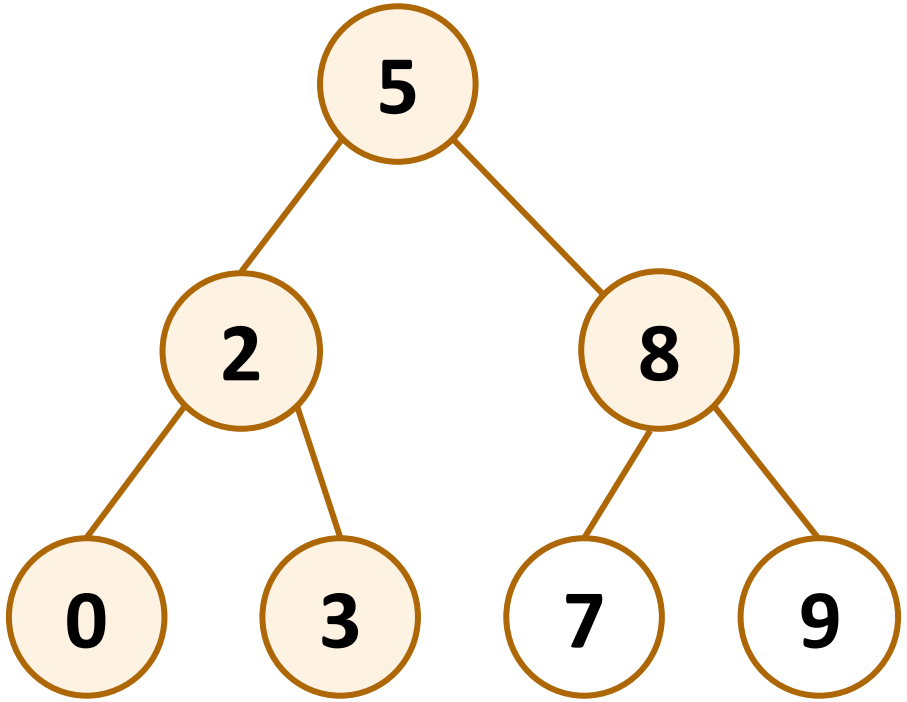


DFS walkthrough

Visited so far



Backtrack to **5** and then continue with its next child **8**.

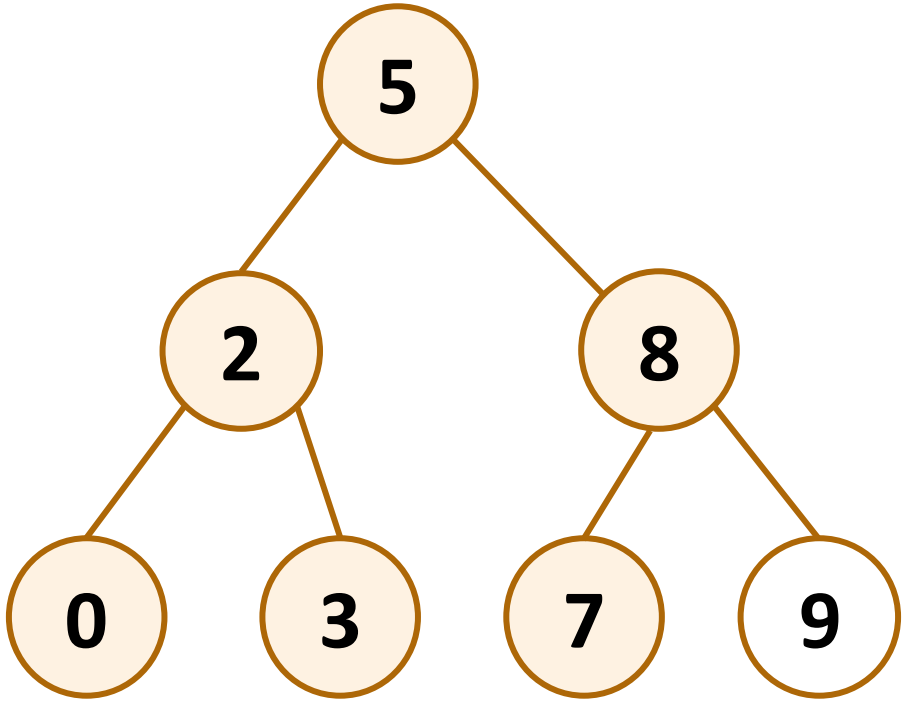


DFS walkthrough

Visited so far

5	2	0	3	8	7

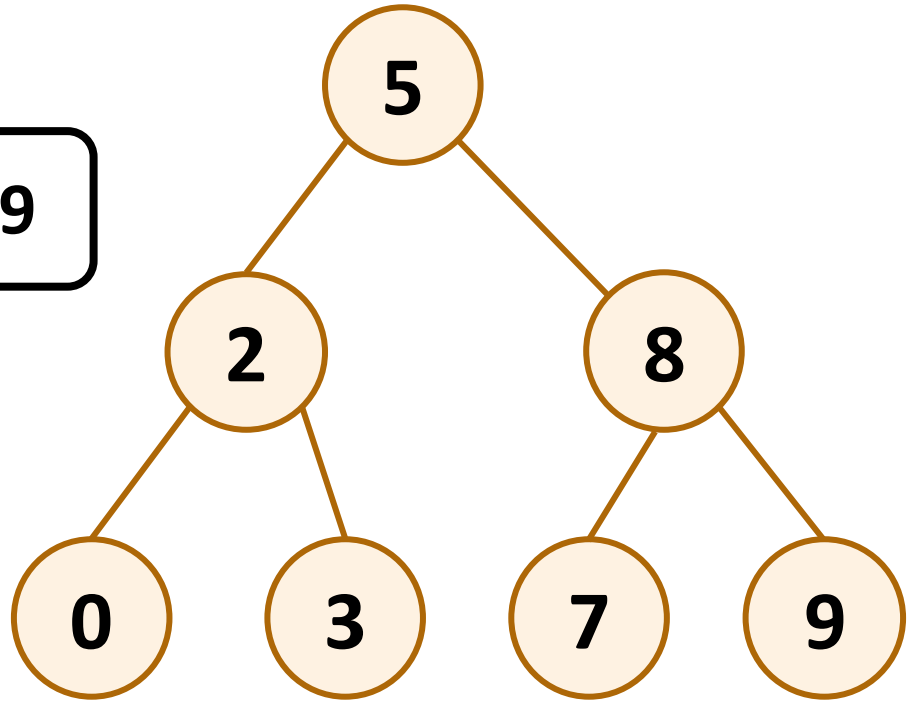
Go to **7**.



DFS walkthrough

Visited so far

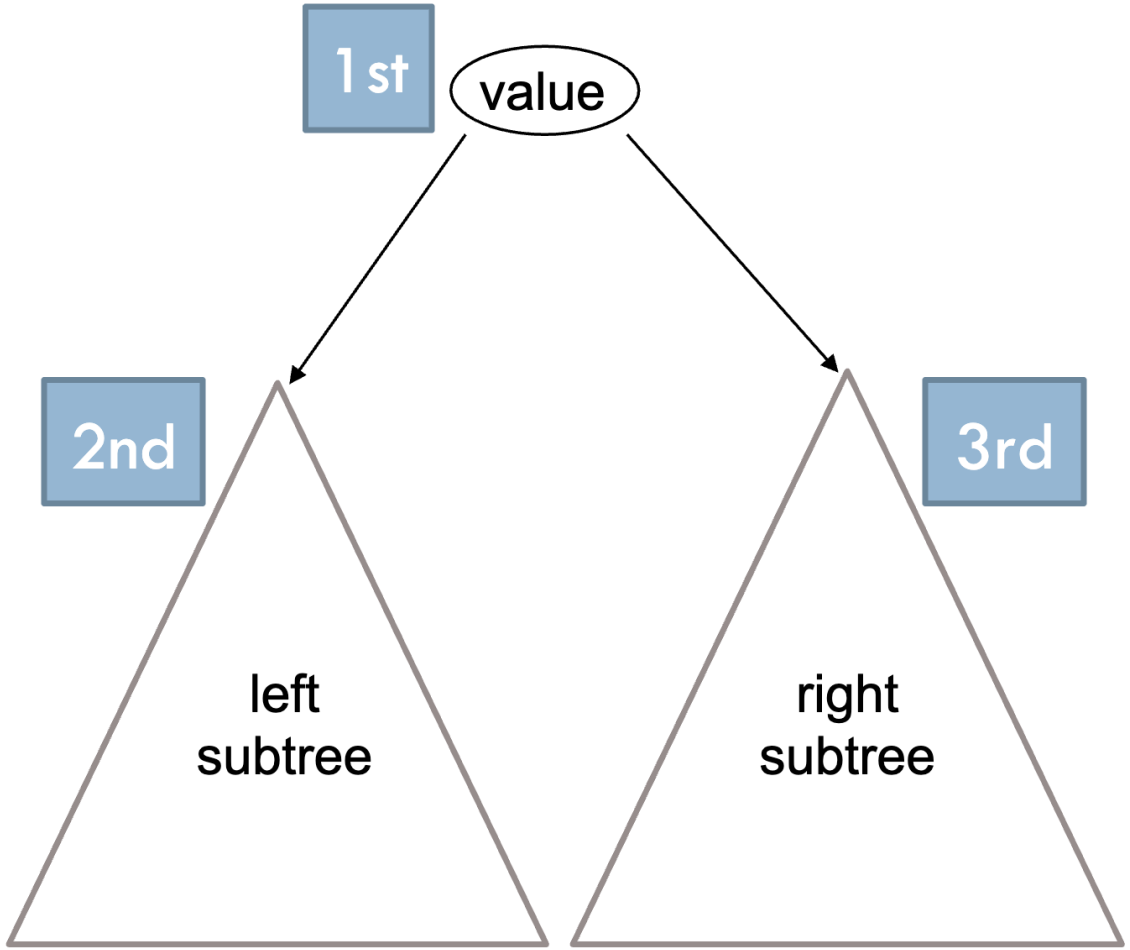
5	2	0	3	8	7	9



Backtrack to **8** and then continue with its next child **9**.

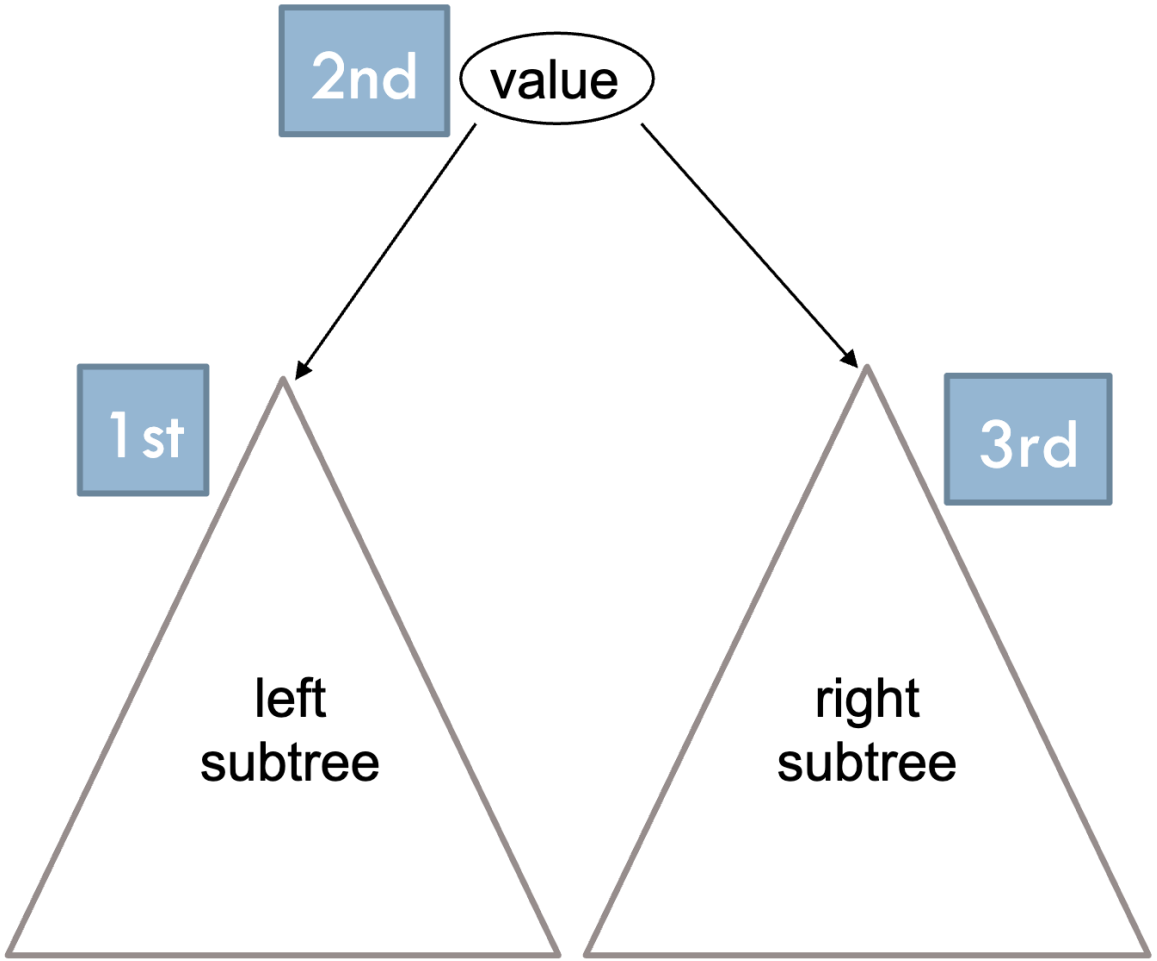
Preorder

Pre: process root before subtrees



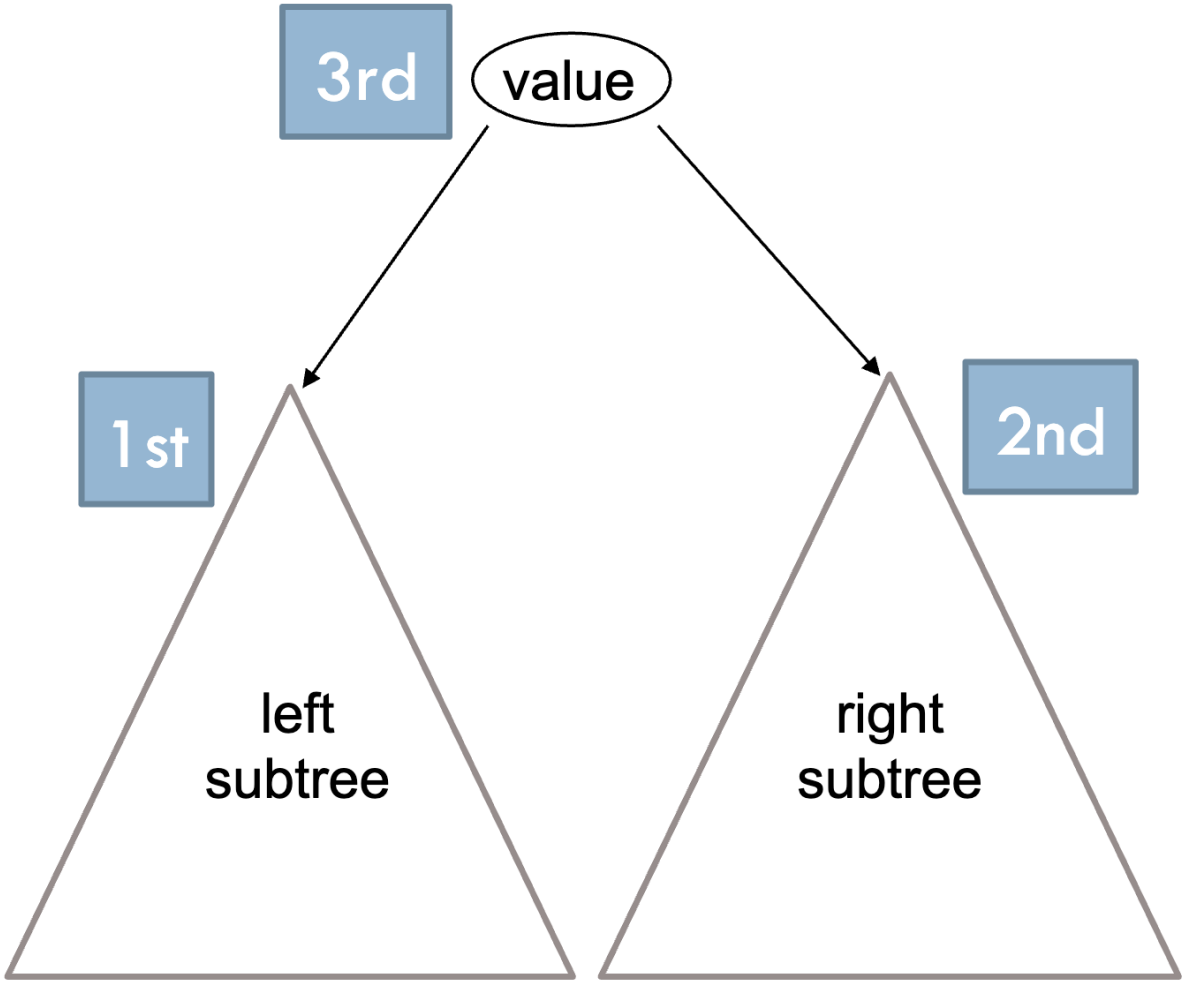
Inorder

In: process root in-between subtrees



Postorder

Post: process root after subtrees

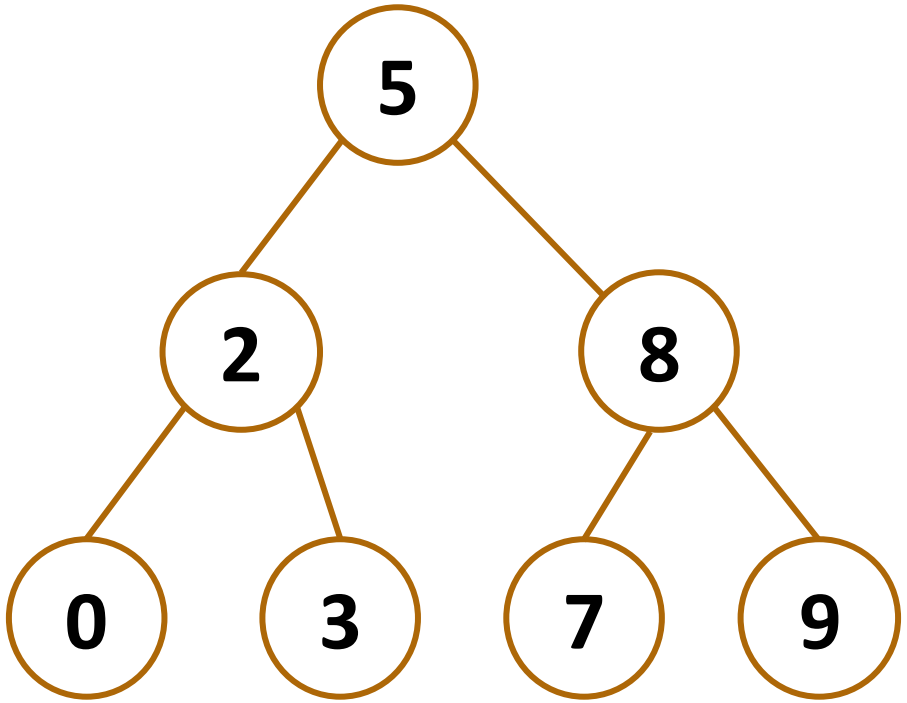


Preorder, inorder, and postorder

Preorder: 5, 2, 0, 3, 8, 7, 9

Inorder: 0, 2, 3, 5, 7, 8, 9

Postorder: 0, 3, 2, 7, 9, 8, 5

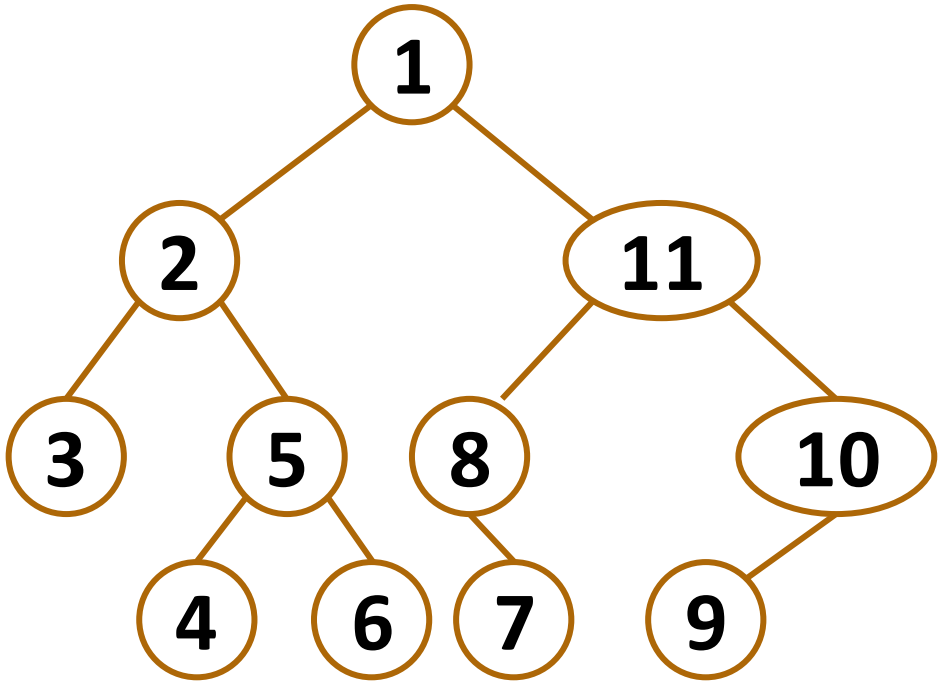


Preorder, inorder, and postorder

Preorder:

Inorder:

Postorder:

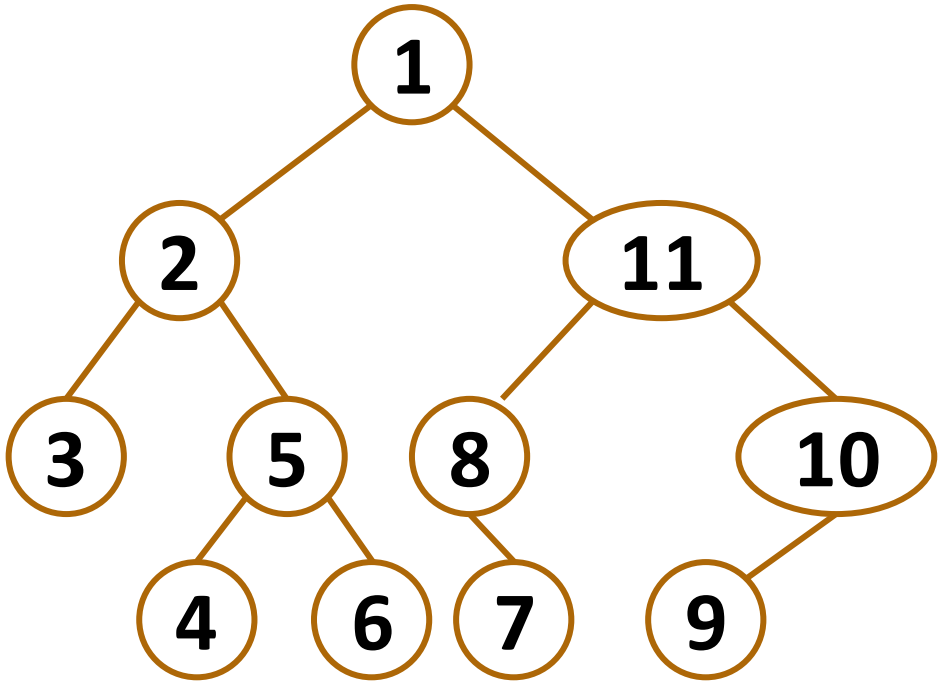


Preorder, inorder, and postorder

Preorder: 1, 2, 3, 5, 4, 6, 11, 8, 7, 10, 9

Inorder: 3, 2, 4, 5, 6, 1, 8, 7, 11, 9, 10

Postorder: 3, 4, 6, 5, 2, 7, 8, 9, 10, 11, 1

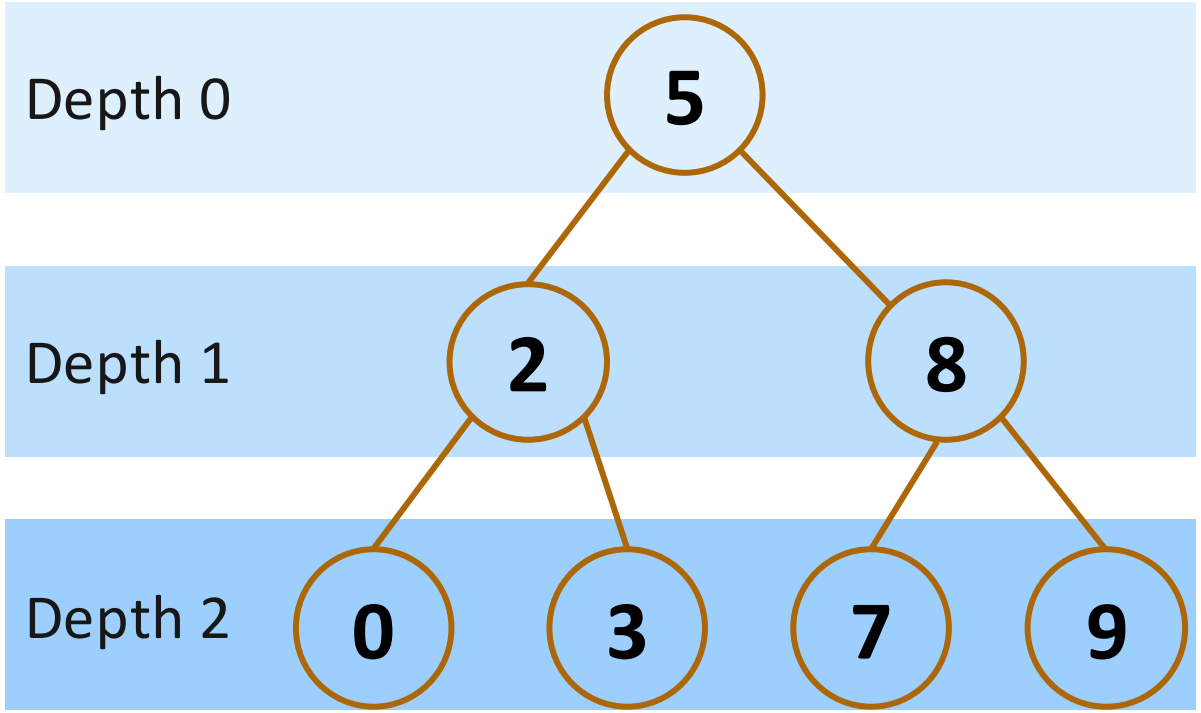


BFS: visit the tree level by level

BFS explores all nodes at depth 0, then depth 1, then depth 2, and so on.

How to implement this?

- uses a queue to maintain the current frontier.

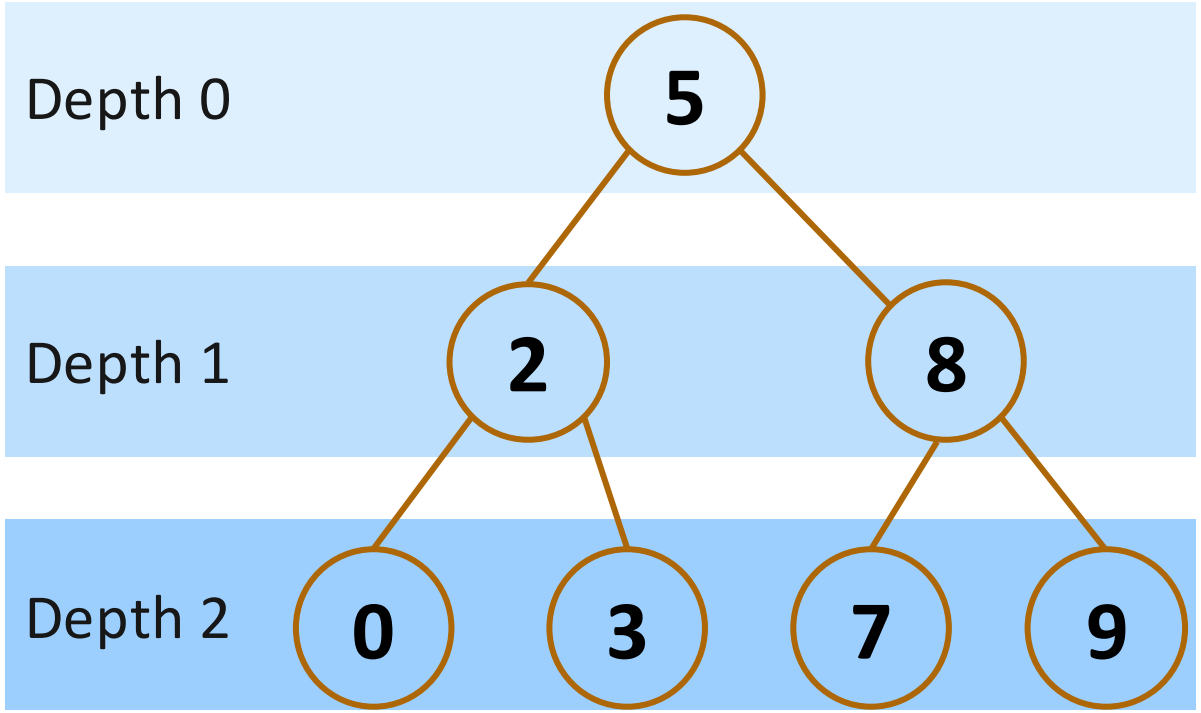


BFS: visit the tree level by level

香港科技大学(广州)
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY (GUANGZHOU)

Recursive idea

```
1 Q ← empty queue
2 enqueue(Q, root)
3 while Q ≠ ∅:
4         u ← dequeue(Q)
5         visit(u)
6         for v in children[u]:
7                 enqueue(Q, v)
```



visit u first

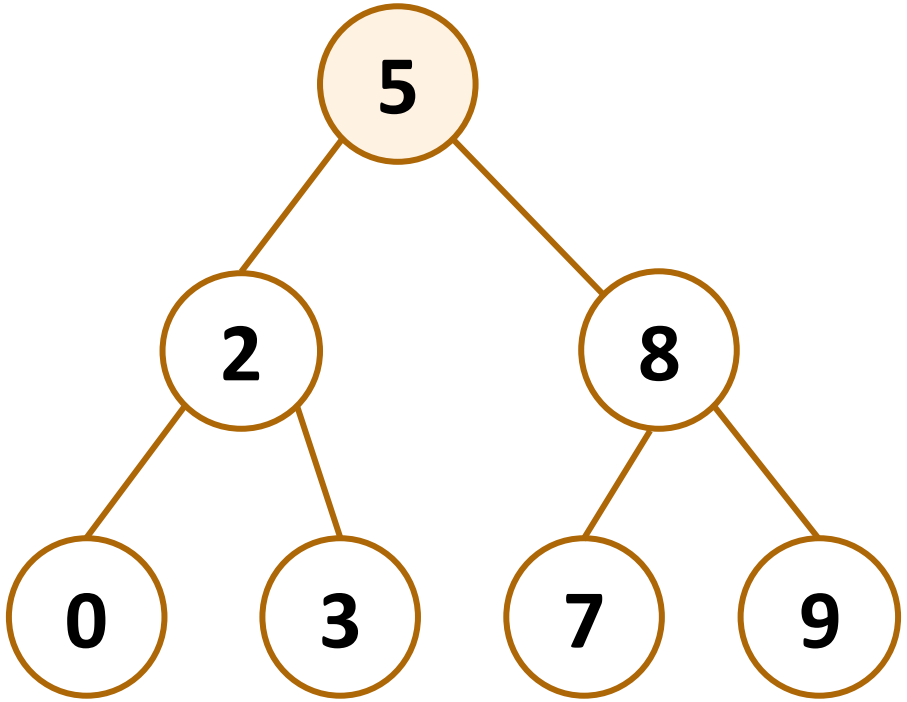
then recursively process every child

BFS walkthrough

Visited so far



All children of the root join the queue in one step.

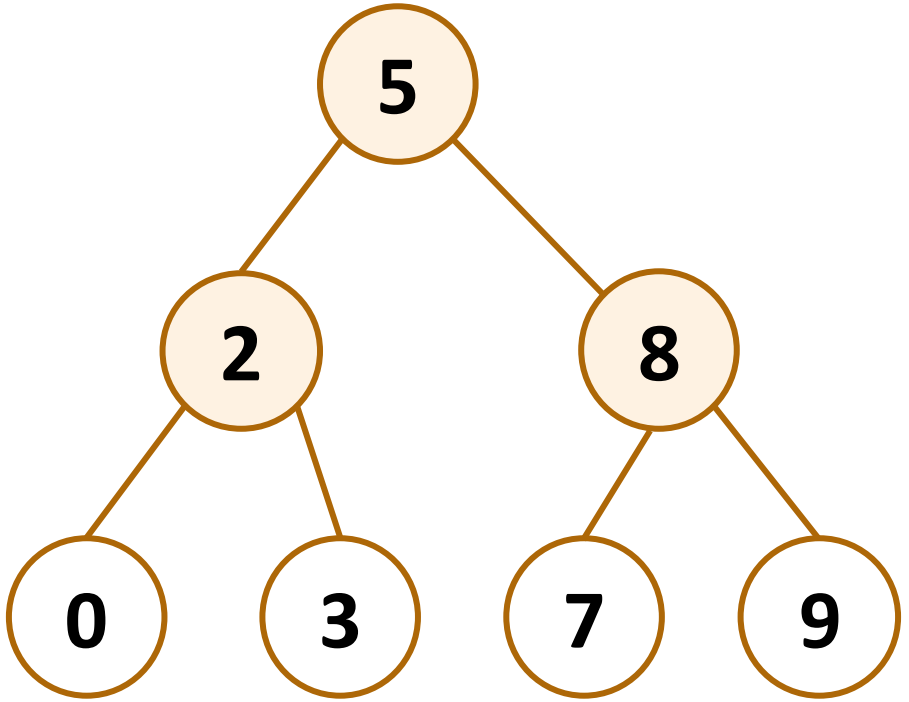


BFS walkthrough

Visited so far



After all depth-1 nodes are processed, the queue contains depth-2 nodes.

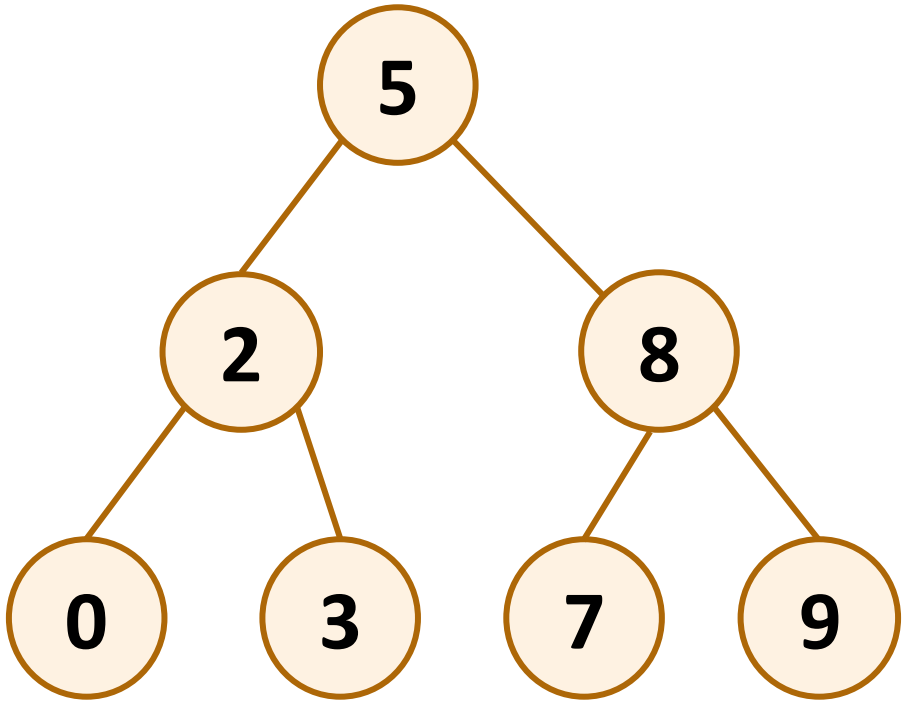


BFS walkthrough

Visited so far

5	2	8	0	3	7	9

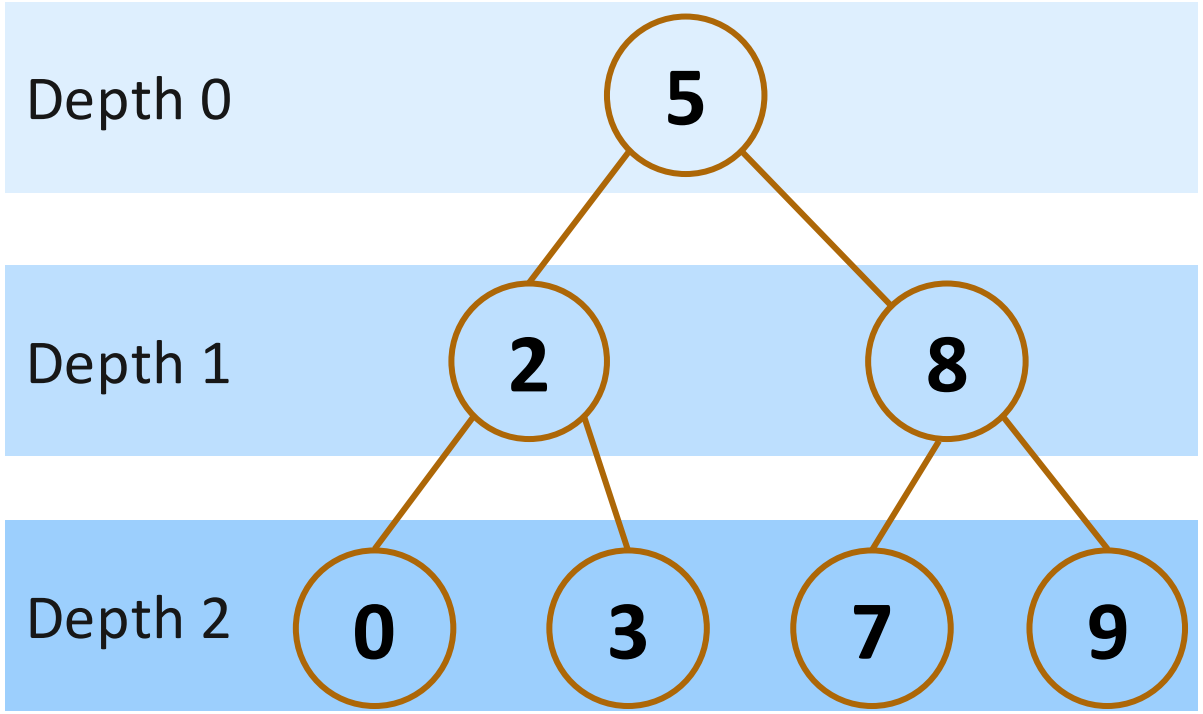
After all depth-1 nodes are processed, the queue contains depth-2 nodes.



BFS immediately gives node depths

BFS visits nodes level by level, the first time we reach a node we know its depth.

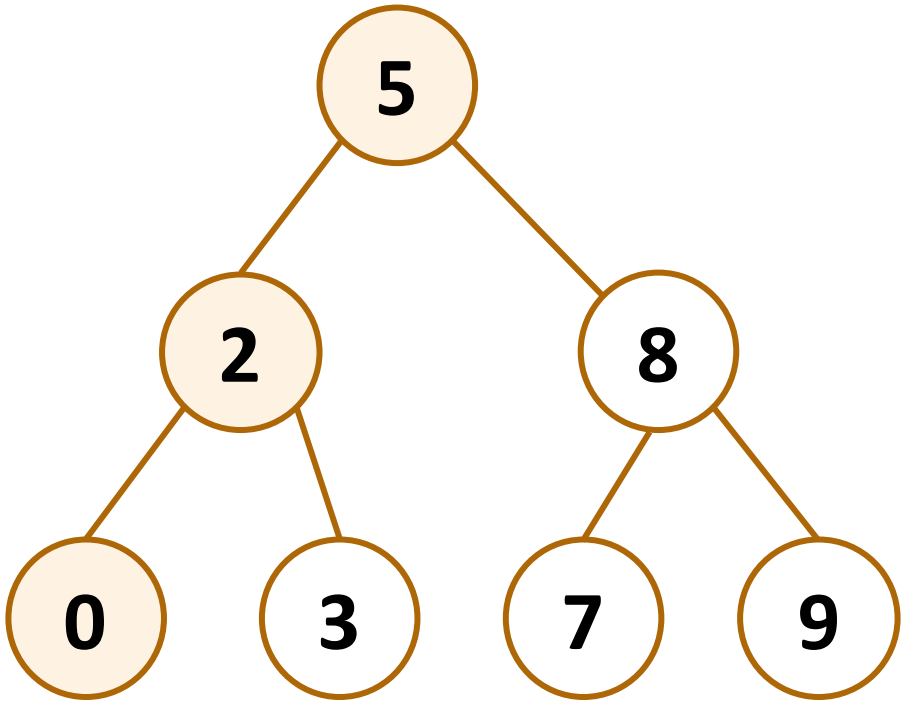
Depth	Nodes
0	5
1	2, 8
2	0, 3, 7, 9



DFS vs BFS

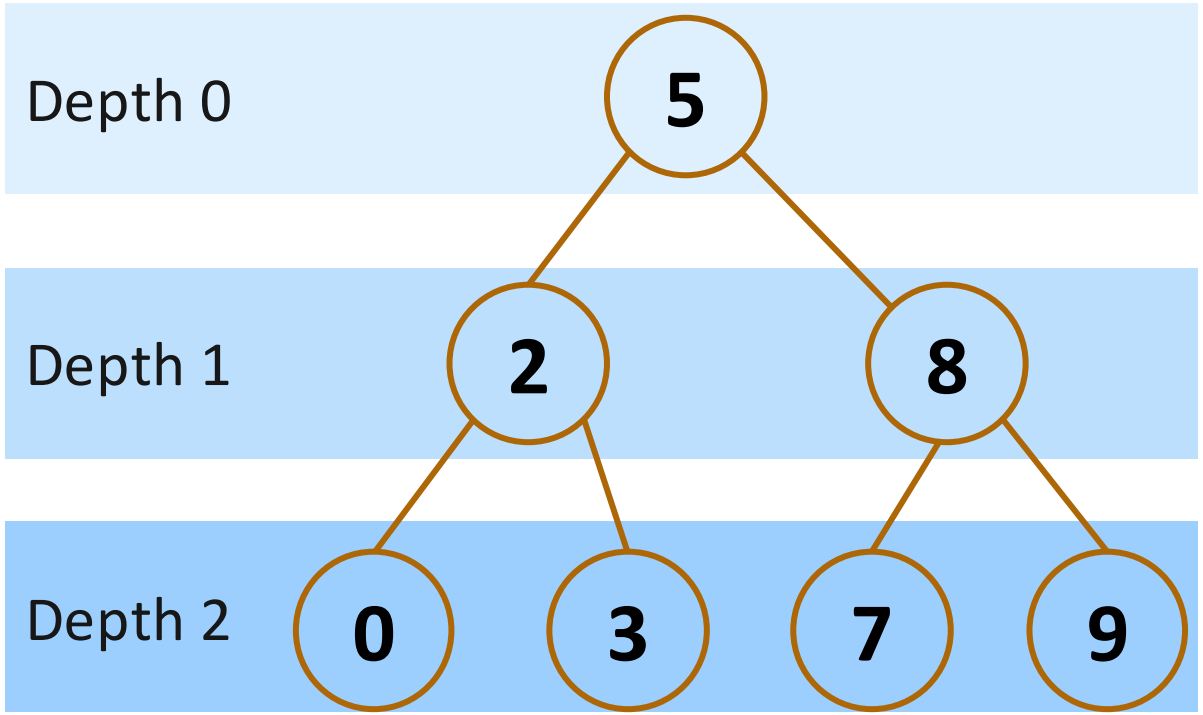
香港科技大学(广州)
THE HONG KONG
UNIVERSITY OF SCIENCE AND
TECHNOLOGY (GUANGZHOU)

DFS



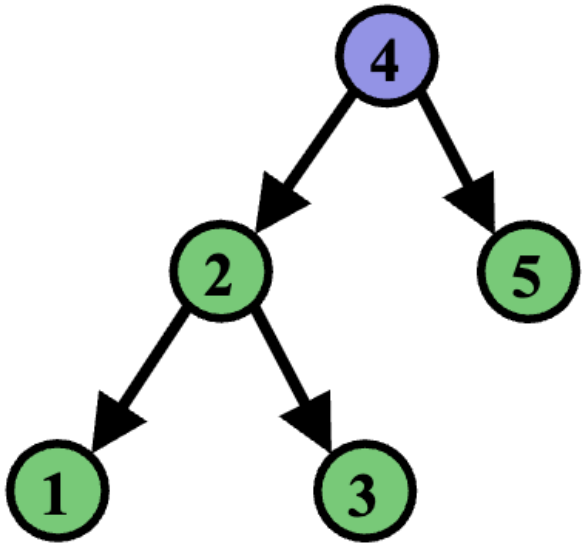
Typical order: 5, 2, 0, 3, 8, 7, 9

BFS

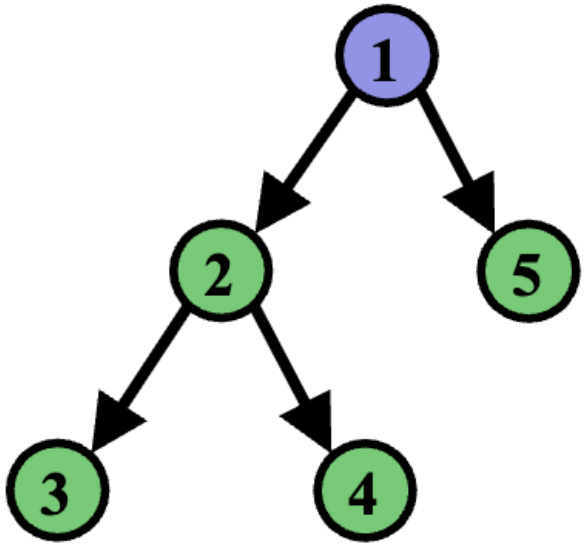


Typical order: 5, 2, 8, 0, 3, 7, 9

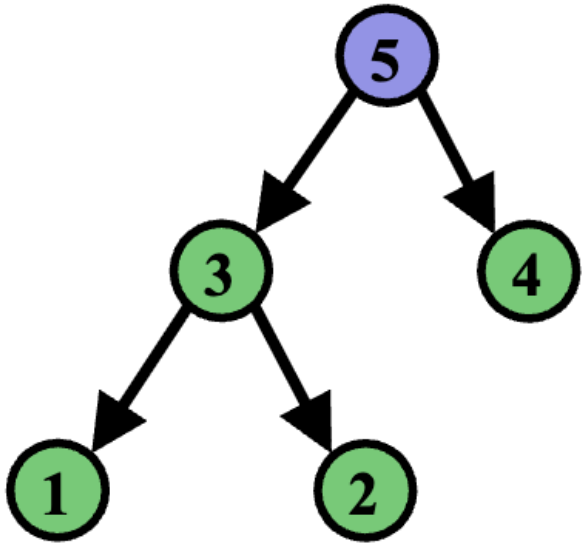
Tree Traversals



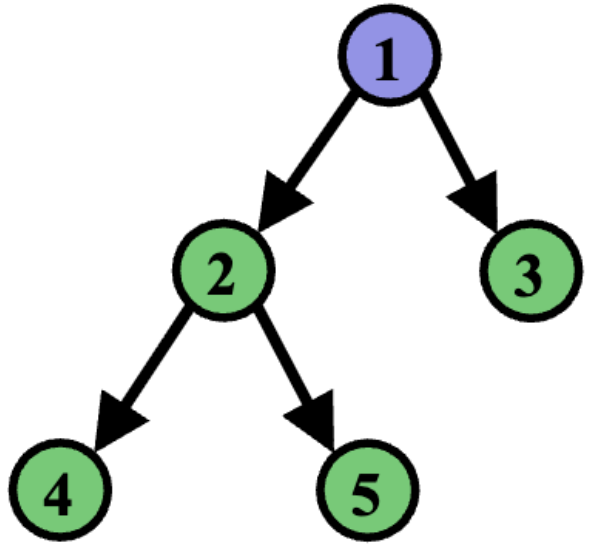
In-order



Pre-order



Post-order



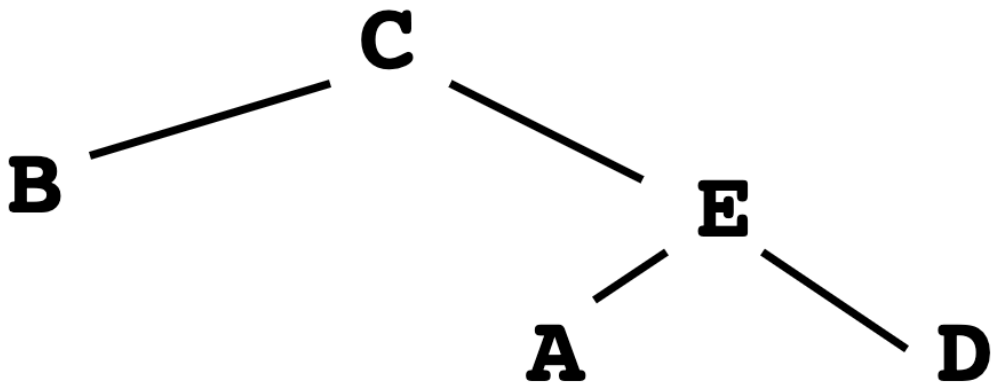
level numbering

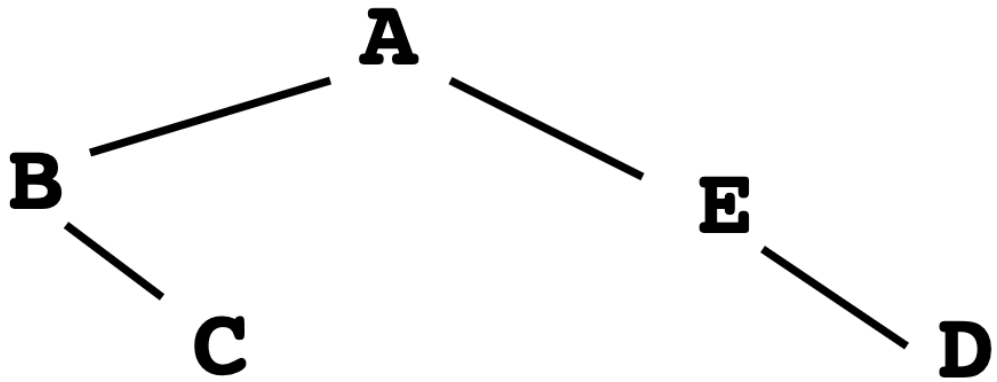
Recover tree from traversal

Suppose inorder is B C A E D.

Can we recover the tree uniquely?

No!





Recover tree from traversal

Suppose inorder is B C A E D

preorder is A B C D E

Can we determine the tree uniquely? Yes!

- What is root? Preorder tells us: A
- What comes before/after root A? Inorder tells us:
  - Before (left subtree of A): B C
  - After (right subtree of A): E D
- Now recurse! Figure out left/right subtrees using same technique.

Recover tree from traversal

香港科技大学(广州)
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY (GUANGZHOU)

Suppose inorder is B C A E D

preorder is A B C D E

Root is A. left subtree contains B C; right contains E D

Left:

Inorder is B C

Preorder is B C

What is root? Preorder tells us: B

- What comes before/after B?
- Inorder tells us:
  - Before: nothing
  - After: C

Right:

Inorder is E D

Preorder is D E

What is root? Preorder tells us: D

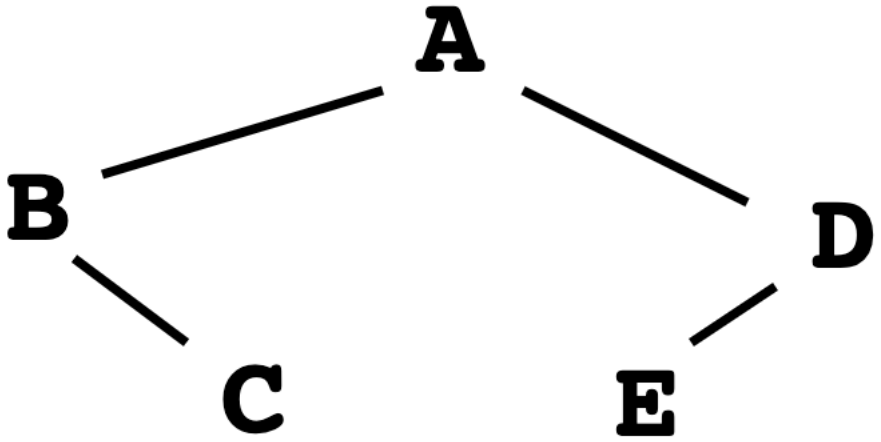
- What comes before/after D?
- Inorder tells us:
  - Before: E
  - After: nothing

Recover tree from traversal

香港科技大学(广州)
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY (GUANGZHOU)

Suppose inorder is B C A E D
preorder is A B C D E

Tree is



Recover tree from traversal

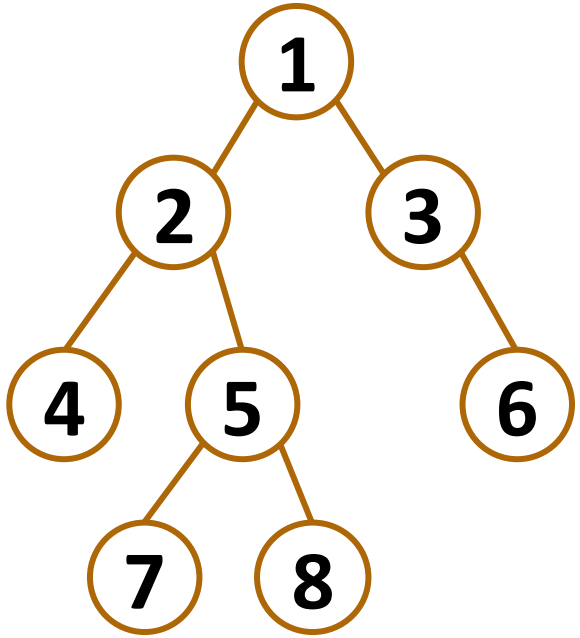
A binary tree can be uniquely determined by

- **(pre-order, in-order)**
- **(post-order, in-order)**

Preorder: 1, 2, 4, 5, 7, 8, 3, 6

Inorder: 4, 2, 7, 5, 8, 1, 3, 6

Postorder: 4, 7, 8, 5, 2, 6, 3, 1



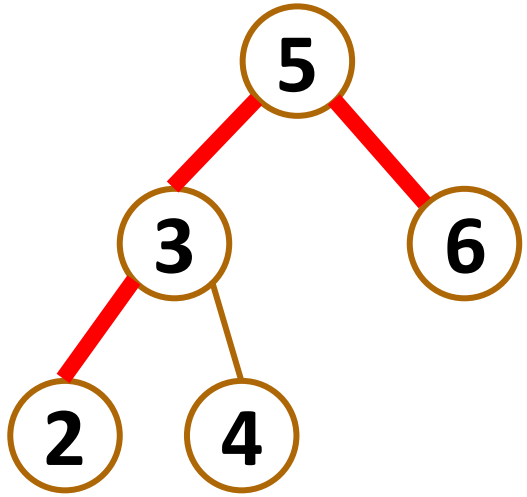
Dynamic Programming on Trees

Diameter of a tree

Definition: the length of the longest path between any two nodes of the tree.

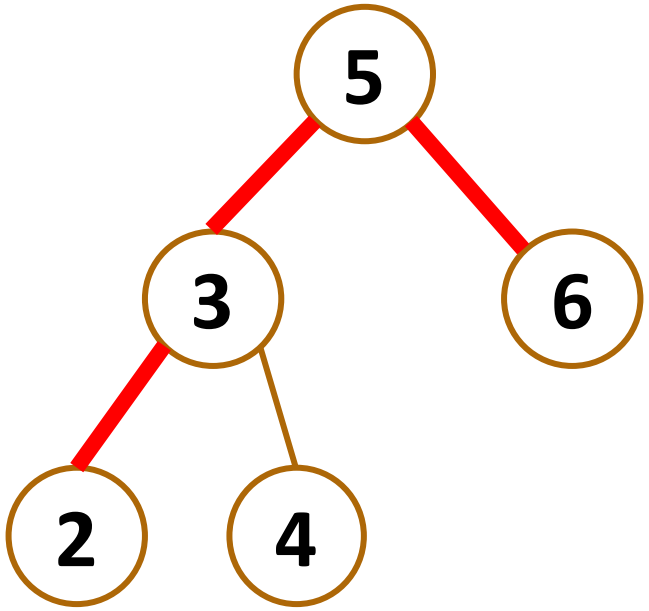
We count edges for the diameter.

- There may be many long paths, but only one maximum length value.

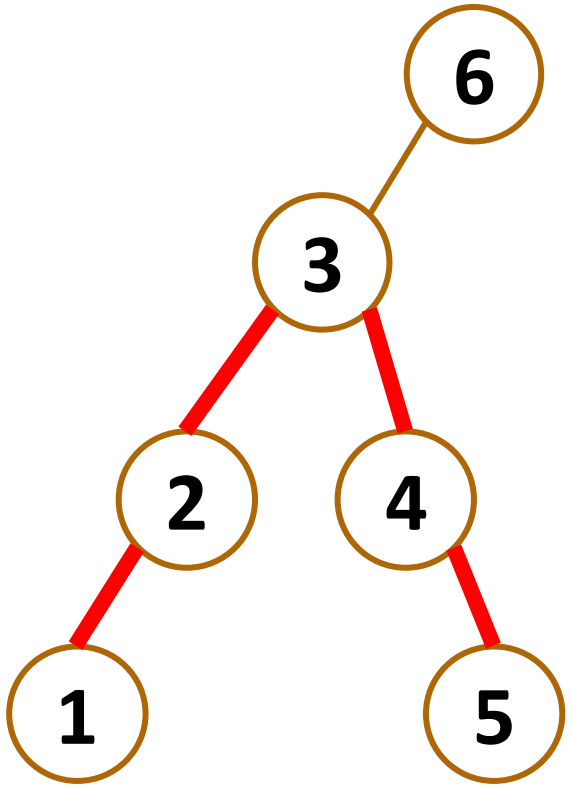


the diameter is 3

Diameter may not pass through the root



diameter is 3



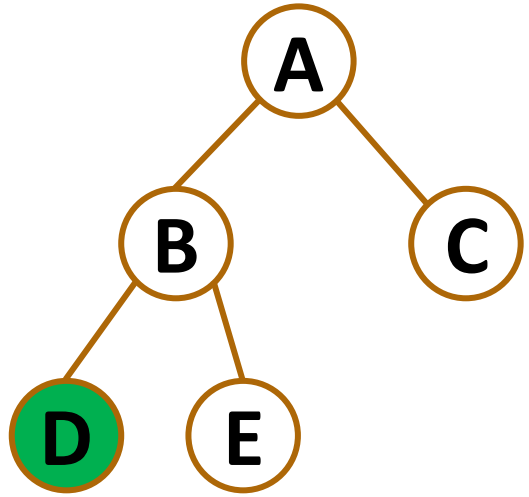
diameter is 4

Naive method

Run **DFS/BFS** from every node once and take the maximum distance from each run, resulting in an algorithm with complexity $O(N^2)$, where N is the number of nodes.

Step 1: choose **D** as a start

Step 2:
dist(D, D) = 0
dist(D, B) = 1
dist(D, E) = 2
dist(D, A) = 2
dist(D, C) = 3

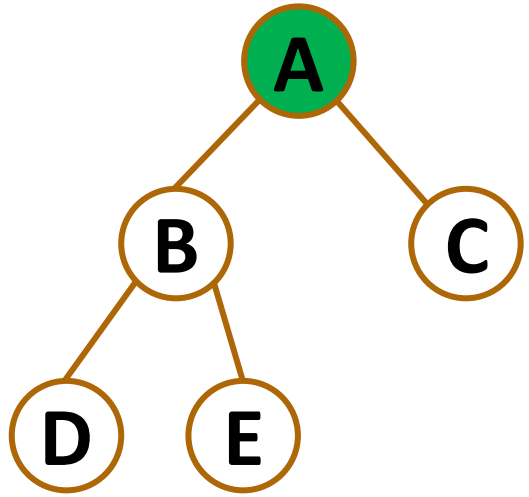


Naive method

Run DFS from every node once and take the maximum distance from each run, resulting in an algorithm with complexity $O(N^2)$, where N is the number of nodes.

Step 3: choose **A** as a start

Step 4:
dist(A, A) = 0
dist(A, B) = 1
dist(A, C) = 1
dist(A, D) = 2
dist(A, E) = 2

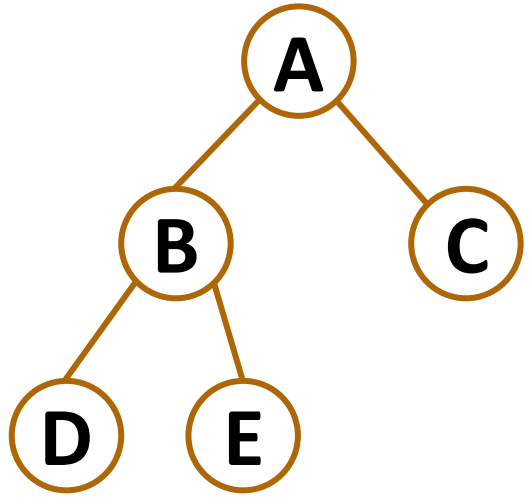


Naive method

Run DFS from every node once and take the maximum distance from each run, resulting in an algorithm with complexity $O(N^2)$, where N is the number of nodes.

Step 5: repeat steps for node B, C, E

Step 6: choose the maximum value across all distances.



The diameter of this tree is 3

2-DFS/BFS

We do not have to run a DFS/BFS from every node: we only need to do it **twice**.

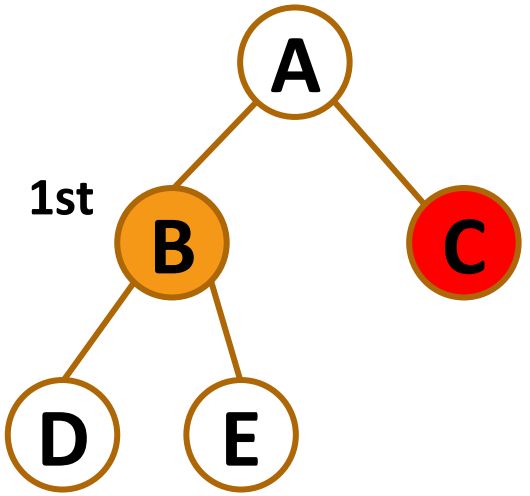
- First, we run a DFS/BFS from any node, recording the depth of each node from the starting node.
- Among all the nodes with the maximum depth, choose any of them, and run another DFS/BFS while recording the depth of each node.
- The maximum depth in the second DFS/BFS is the tree diameter.

This results in an algorithm with time complexity O(N), where N is the number of nodes.

- First, we run a DFS/BFS from any node, recording the depth of each node from the starting node.

We randomly choose node **B:**

dist(B, B) = 0
dist(B, A) = 1
dist(B, D) = 1
dist(B, E) = 1
dist(B, C) = 2



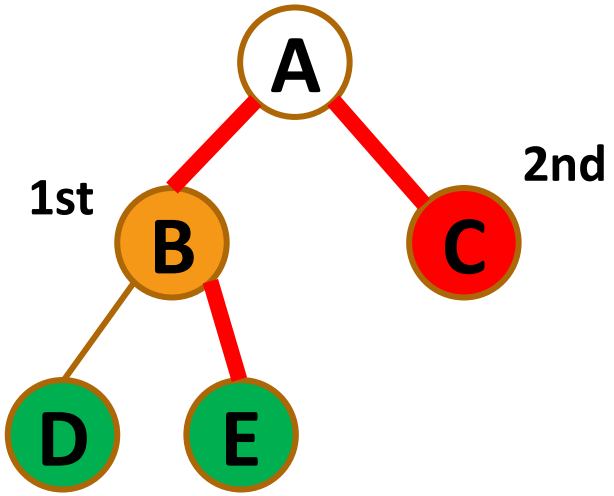
The maximum depth for node B is 2

2-DFS/BFS

- Among all the nodes with the maximum depth, choose any of them, and run another DFS/BFS while recording the depth of each node.

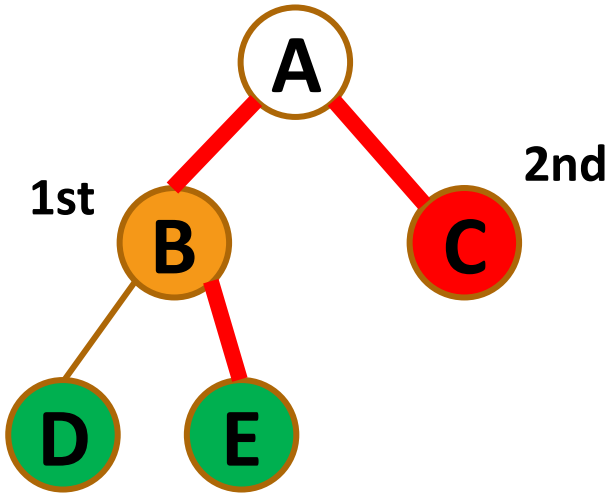
We choose node **C:**

dist(C, C) = 0
dist(C, A) = 1
dist(C, B) = 2
dist(C, D) = 3
dist(C, E) = 3



The maximum depth for node C is 3, and the diameter of this tree is 3

```
Global max_dist = 0
void dfs(int x, int parent, int dist){
    if (dist >= max_dist){
        deepest_node = x;
        max_dist = dist;
    }
    for (auto i : e[x])
        if (i != parent) dfs(i, x, dist + 1);
}
int tree_diameter(){
    max_dist = 0;
    dfs(1, -1, 0);
    max_dist = 0;
    dfs(deepest_node, -1, 0);
    return max_dist;
}
```



Why does it work?

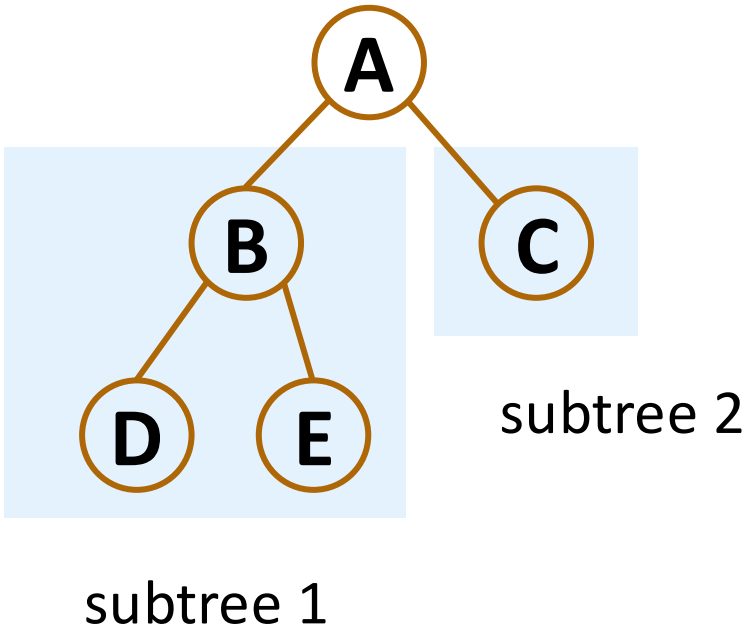
Key observation for diameter

香港科技大学(广州)
THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY (GUANGZHOU)

At a node u, the best path is one of two types:

- It is entirely contained inside one child subtree.
-
- Or it passes through u and uses the two deepest child branches.

That suggests we should keep the **best** downward path from every node.

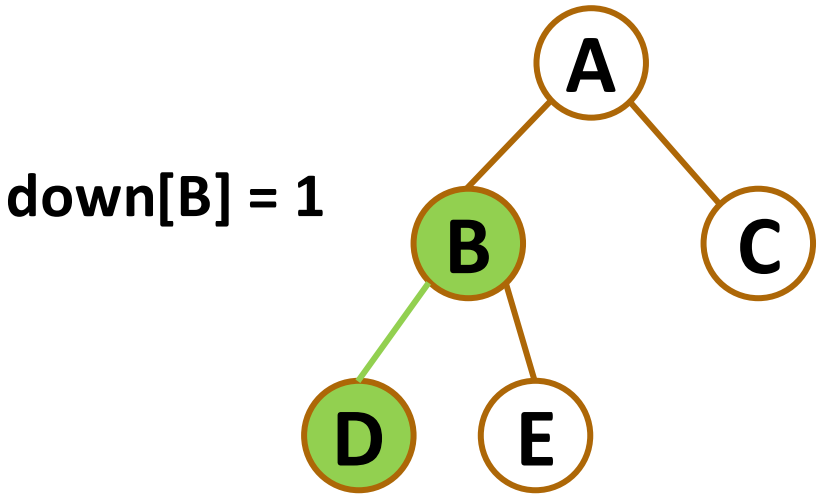


Tree DP

- Assume a tree is rooted
- Use nodes as DP states
- Use parent-children to define transition formula

Tree DP - State

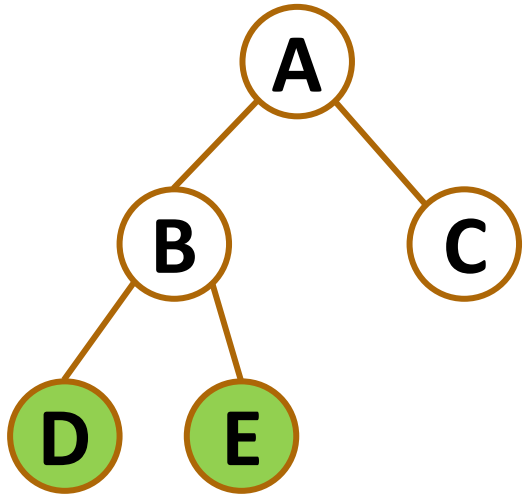
down[u] = the length of the longest path that starts at u and goes downward into its subtree.



Tree DP - Base case

- Leaf: down[u] = 0

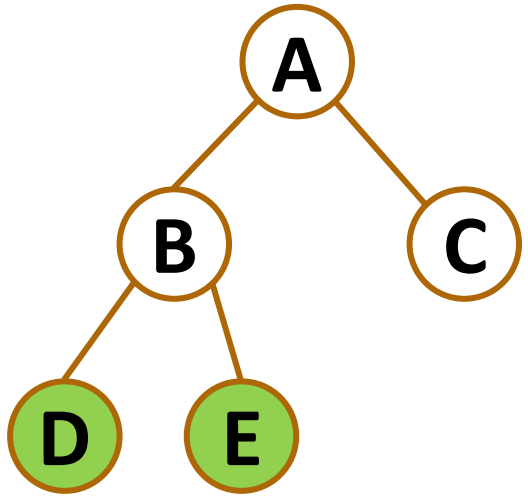
No edge can be taken downward.



Tree DP - Transition Formula

Suppose u has children v_1, v_2, …, v_k
For each child v, if we go from u to v, then continue along the best downward path of v, the length becomes:

$$down[v] + 1$$



Let
best1 = the largest value among all *down[v] + 1*
best2 = the second largest value among all *down[v] + 1*
(best1 and best2 are associated with different v_i)

Tree DP - Transition Formula

Longest downward path from u

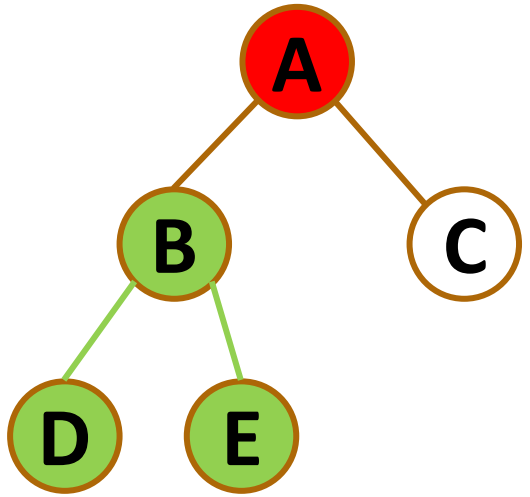
$$down[u] = best1$$

Longest path passing through u

$$through[u] = best1 + best2$$

The diameter is the maximum such value over all nodes:

$$diameter = \max through[u]$$



through[B] = 1 + 1 = 2

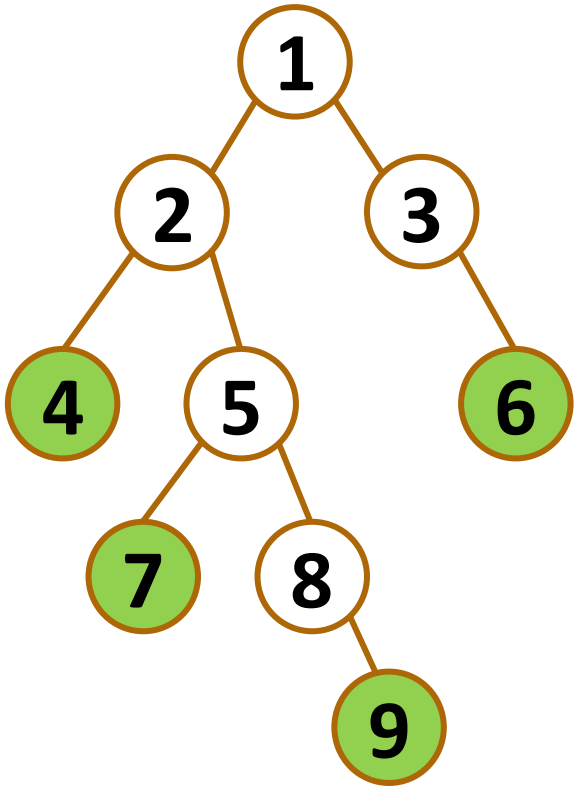
through[A] = 2 + 1 = 3

diameter = through[A] = 3

Tree DP - another example

For every leaf u,
down[u] = 0

down[4] = down[7] = down[9] =
down[6] = 0



Tree DP - another example

down[3] = 1 + down[6] = 1

best1 = 1, best2 = 0

down[8] = 1 + down[9] = 1

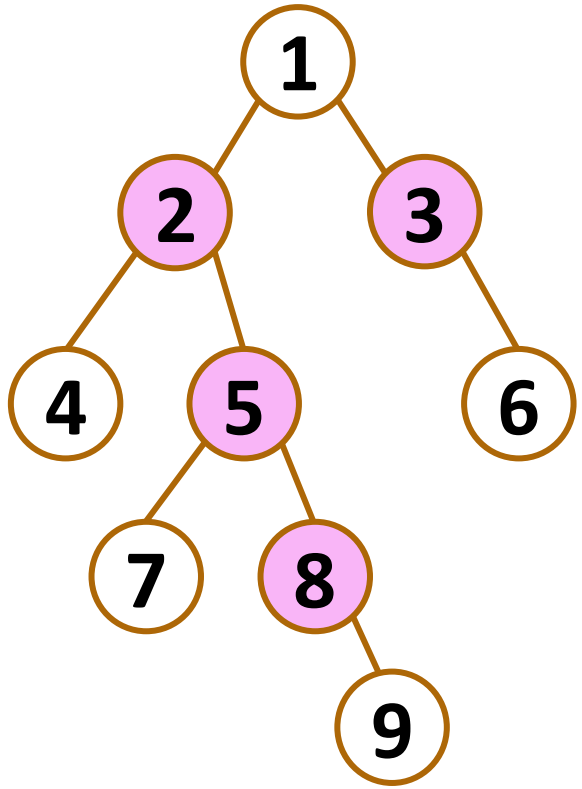
best1 = 1, best2 = 0

down[5] = 1 + max(down[7] +
down[8]) = 1 + max(0, 1) = 2

best1 = 2, best2 = 1

down[2] = 1 + max(down[4] +
down[5]) = 3

best1 = 3, best2 = 1

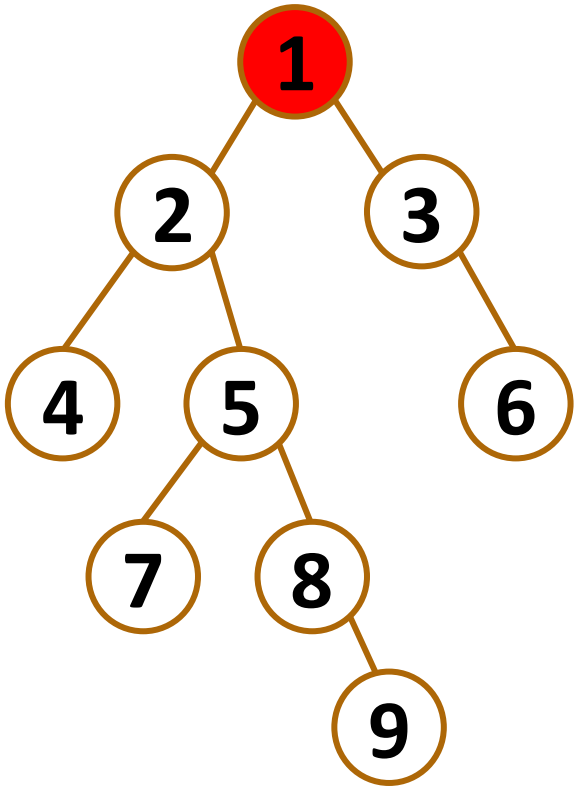


Tree DP - another example

down[1] = 1 + max(down[2] +
down[3]) = 4

best1 = 4, best2 = 2

the diameter = best1 + best2 = 6



Tree DP

```
1 ans = 0
2
3 def dfs(u, parent):
4    global ans
5    best1 = best2 = 0
6    for v in adj[u]:
7        if v == parent: continue
8        down_v = 1 + dfs(v, u)
9        if down_v > best1:
10           best2 = best1
11           best1 = down_v
12       elif down_v > best2:
13           best2 = down_v
14    ans = max(ans, best1 + best2)
```

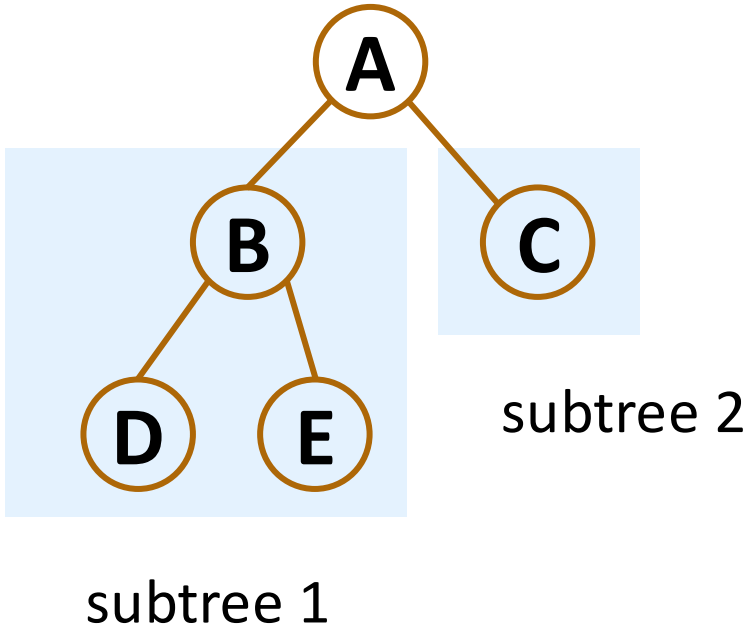
- **Time complexity: O(N)**

Tree DP - Subtree size

- Given a rooted tree of size N
- Calculate the size of each subtree

- dp[leaf] = 1
- dp[u] = size of subtree u
- dp[u] = 1 + sum(dp[v]) where v is u's children

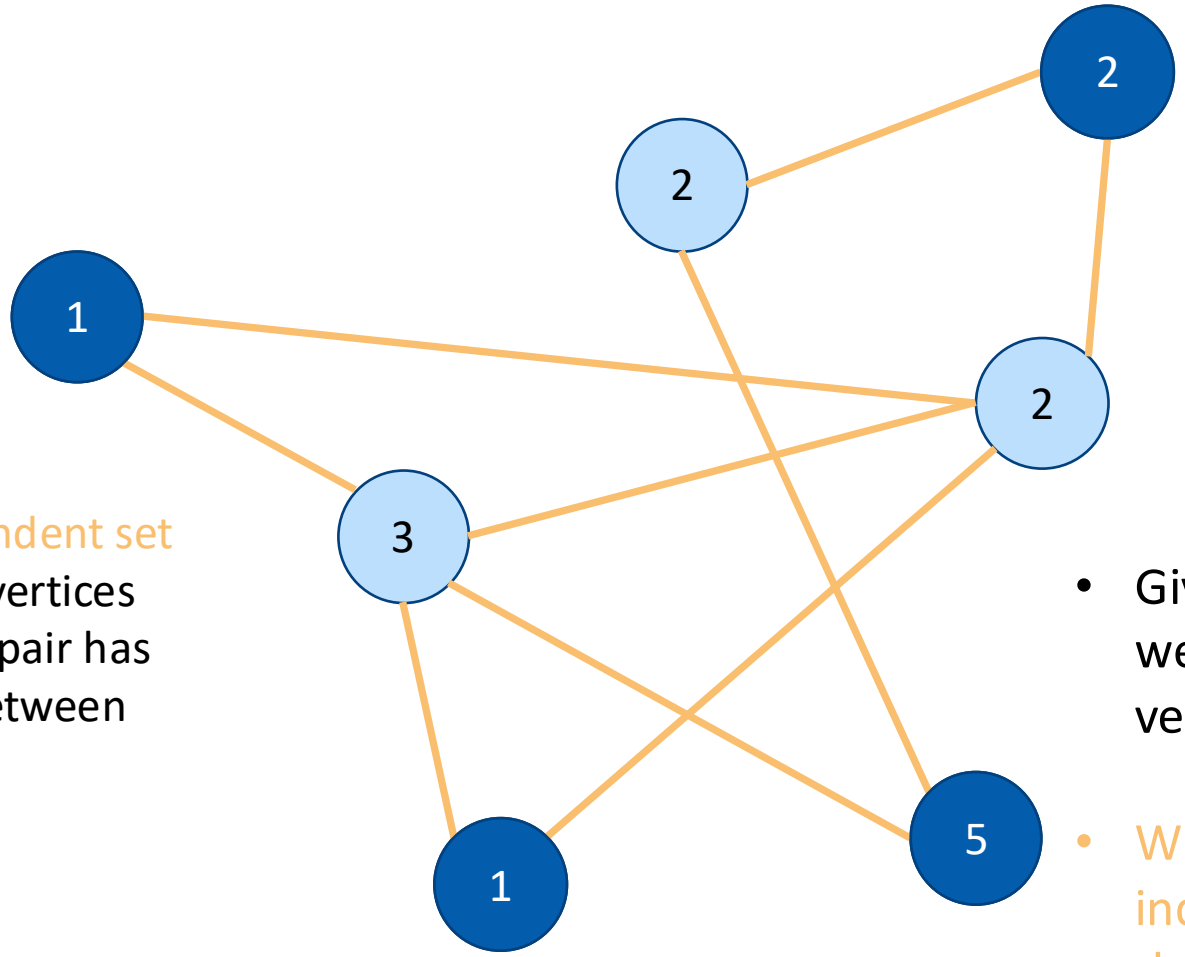
- Each node is visited once only
- Time complexity = O(N)



Tree DP - Subtree size


```
Algorithm ComputeSize(u)
    size[u] = 1
    for each child v of u do
        ComputeSize(v)
        size[u] = size[u] + size[v]
    return size[u]
```

Independent Set



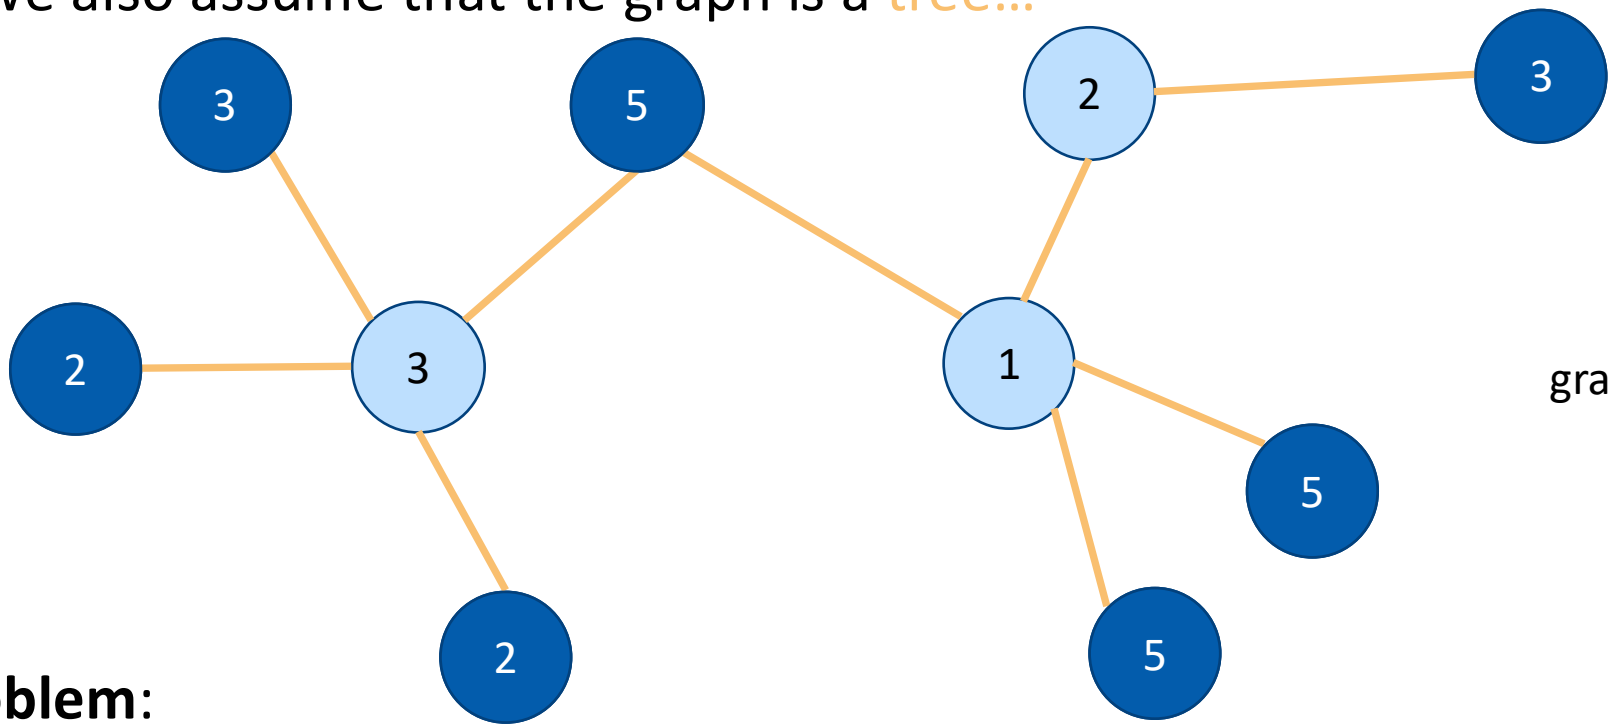
An independent set is a set of vertices so that no pair has an edge between them.



- Given a graph with weights on the vertices…
- What is the independent set with the largest weight?

Independent Set

- Actually, this problem is NP-complete.
  So, we are unlikely to find an efficient algorithm.
- But if we also assume that the graph is a tree…



A tree is a connected graph with no cycles.



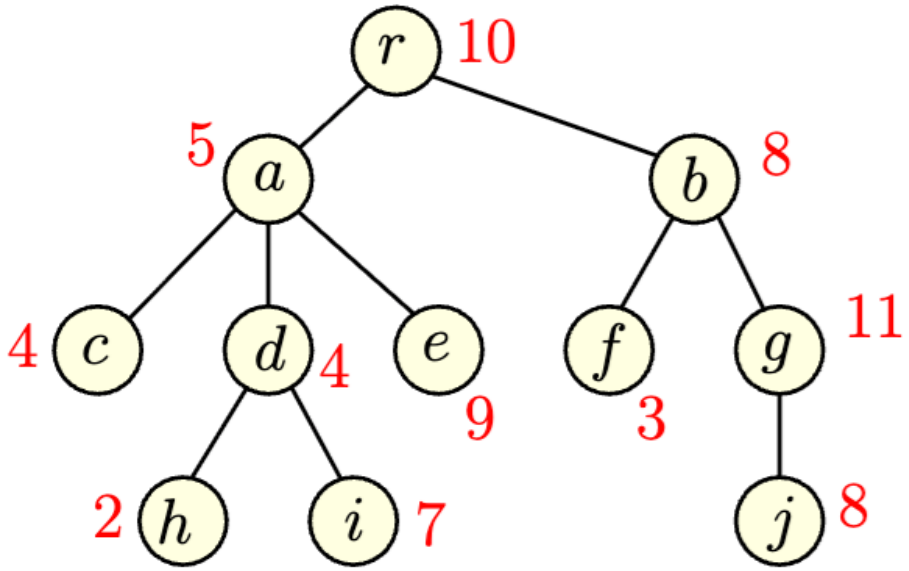
Problem:

find a maximal independent set in a tree (with vertex weights).

Maximum Weight Independent Set in a Tree

Input Tree T = (V, E) and weights w (v) ≥ 0 for each v ∈ V

Goal Find maximum weight independent set in T



Note: We've studied this part in dynamic programming class.

Recipe for applying Dynamic Programming

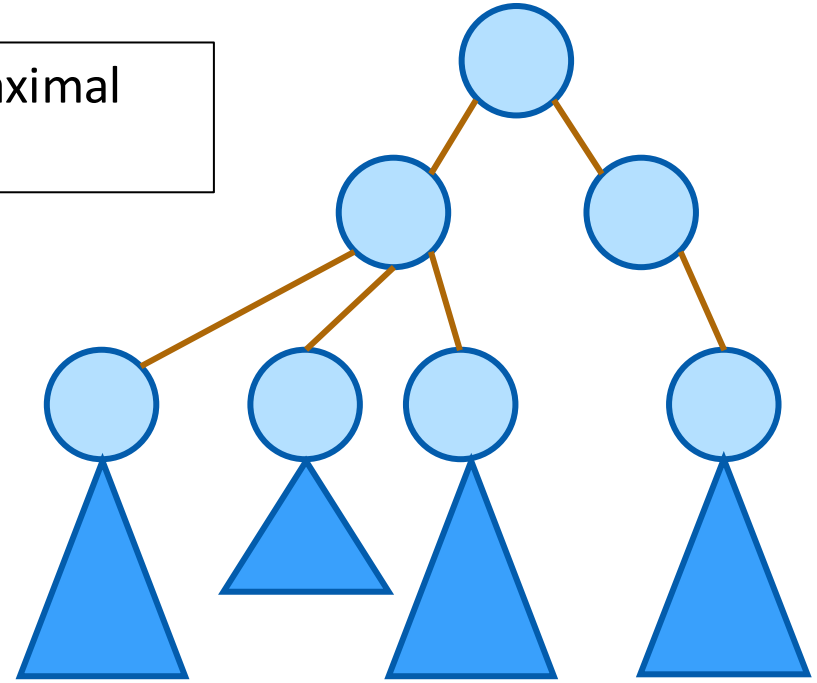
- **Step 1:** Characterize the structure of an optimal solution.
- **Step 2:** Define the optimal value recursively.
- **Step 3:** Compute the optimal values using dynamic programming.
- **Step 4:** If needed, keep track of some additional info so that the algorithm from Step 3 can find the optimal solution.

DP Recipe - Step 1

- Natural candidate for vn is root r of T ? Let O be an optimum solution to the whole problem.

- There are two cases:
  - Case r $\notin$ O : Then O contains an optimum solution for each subtree of T hanging at a child of r.
  - Case r $\in$ O : None of the children of r can be in O. O − {r}contains an optimum solution for each subtree of Thanging at a grandchild of r .

The root is **not** in a maximal independent set.

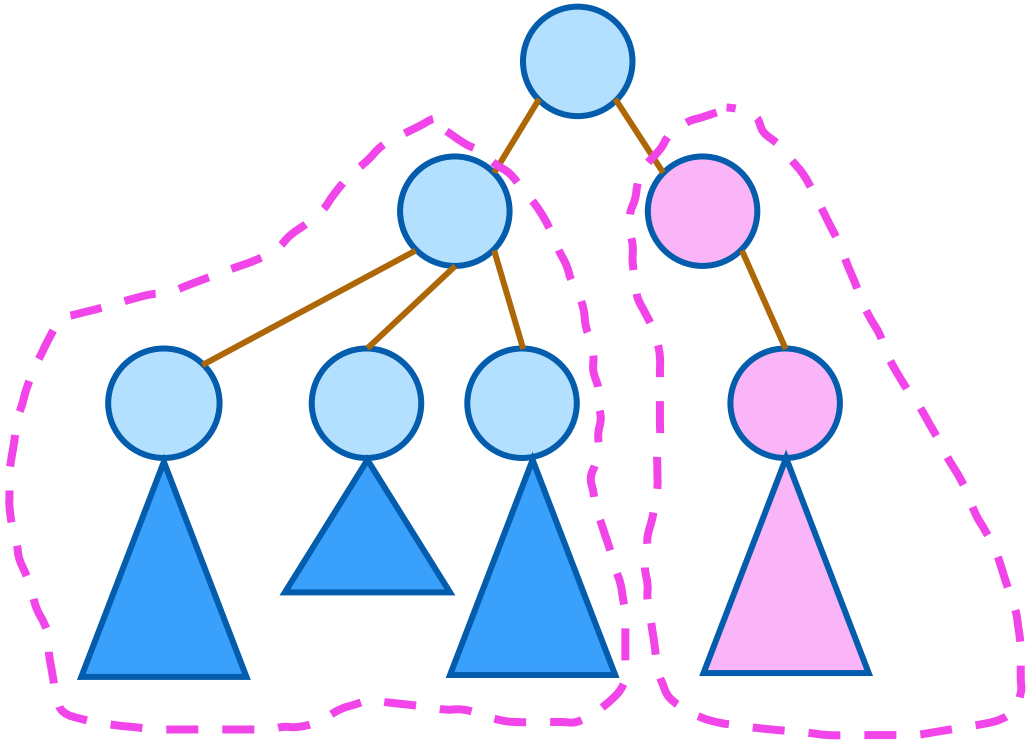


The root is in a maximal independent set.

DP Recipe - Step 1

- Case $r \notin O$: Then O contains an optimum solution for each subtree of T hanging at a child of r.

> The root is **not** in a maximal independent set.

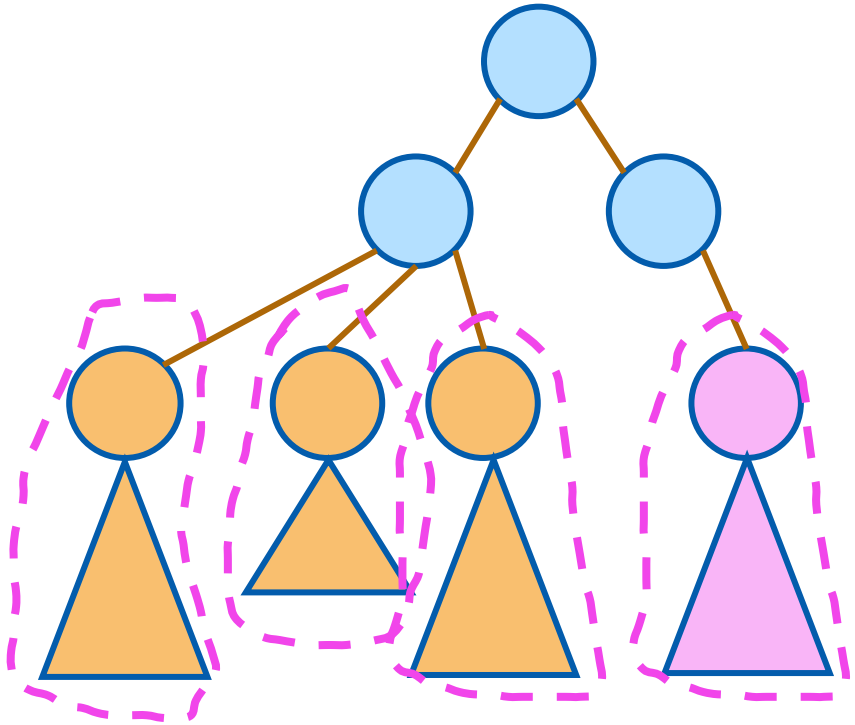


smaller subproblems

DP Recipe - Step 1

- Case $r \in O$: None of the children of r can be in O. O − {r}contains an optimum solution for each subtree of Thanging at a grandchild of r

The root is in a maximal independent set.



smaller subproblems

DP Recipe - Step 2

$T(u)$: subtree of T hanging at node u

$OPT(u)$: max weighted independent set value in $T(u)$

$$OPT(u) = \max \begin{cases} \sum_{v \text{ child of } u} OPT(v), \\ w(u) + \sum_{v \text{ grandchild of } u} OPT(v) \end{cases}$$

